



NoSQL/LevelDB Intro

NoSQL=No SQL or Not Only SQL

- database cannot support big data! What? Are you kidding me?
- for a cluster with 1000 nodes, the CAP cannot be satisfied together?
- CAP=Consistency Availability Partitioning
- How to guarantee that data are 100% available and 100% accessible even hardware fails?
- Conventional database is not designed for this
- NoSQL is a compromise. It sacrifices Consistency for Availability by applying the MVC rule(multi-version concurrency) and eventual consistency
- only support key-based lookup

Popular NoSQL System

Type	Notable examples of this type
Key-Value Cache	Apache Ignite, Coherence, eXtreme Scale, Hazelcast, Infinispan, Memcached, Velocity
Key-Value Store	ArangoDB, Aerospike
Key-Value Store (Eventually-Consistent)	Oracle NoSQL Database, Dynamo, Riak, Voldemort
Key-Value Store (Ordered)	FoundationDB, InfinityDB, LMDB, MemcacheDB
Data-Structures Server	Redis
Tuple Store	Apache River, GigaSpaces
Object Database	Objectivity/DB, Perst, ZopeDB
Document Store	ArangoDB, BaseX, Clusterpoint, Couchbase, CouchDB, DocumentDB, IBM Domino, MarkLogic, MongoDB, Qizx, RethinkDB
Wide Column Store	Amazon DynamoDB, Bigtable, Cassandra, Druid, HBase, Hypertable
Native Multi-model Database	ArangoDB, Cosmos DB, OrientDB

1

What is LevelDB? An open sourced implementation of BigTable by Jeff Dean

LevelDb is designed to process billions of key-value pairs

3

But, We come to the origin: LevelDB

2

- open sourced implementation of bigtable: <https://github.com/google/leveldb>

•

- The NoSQL era starts from LevelDB

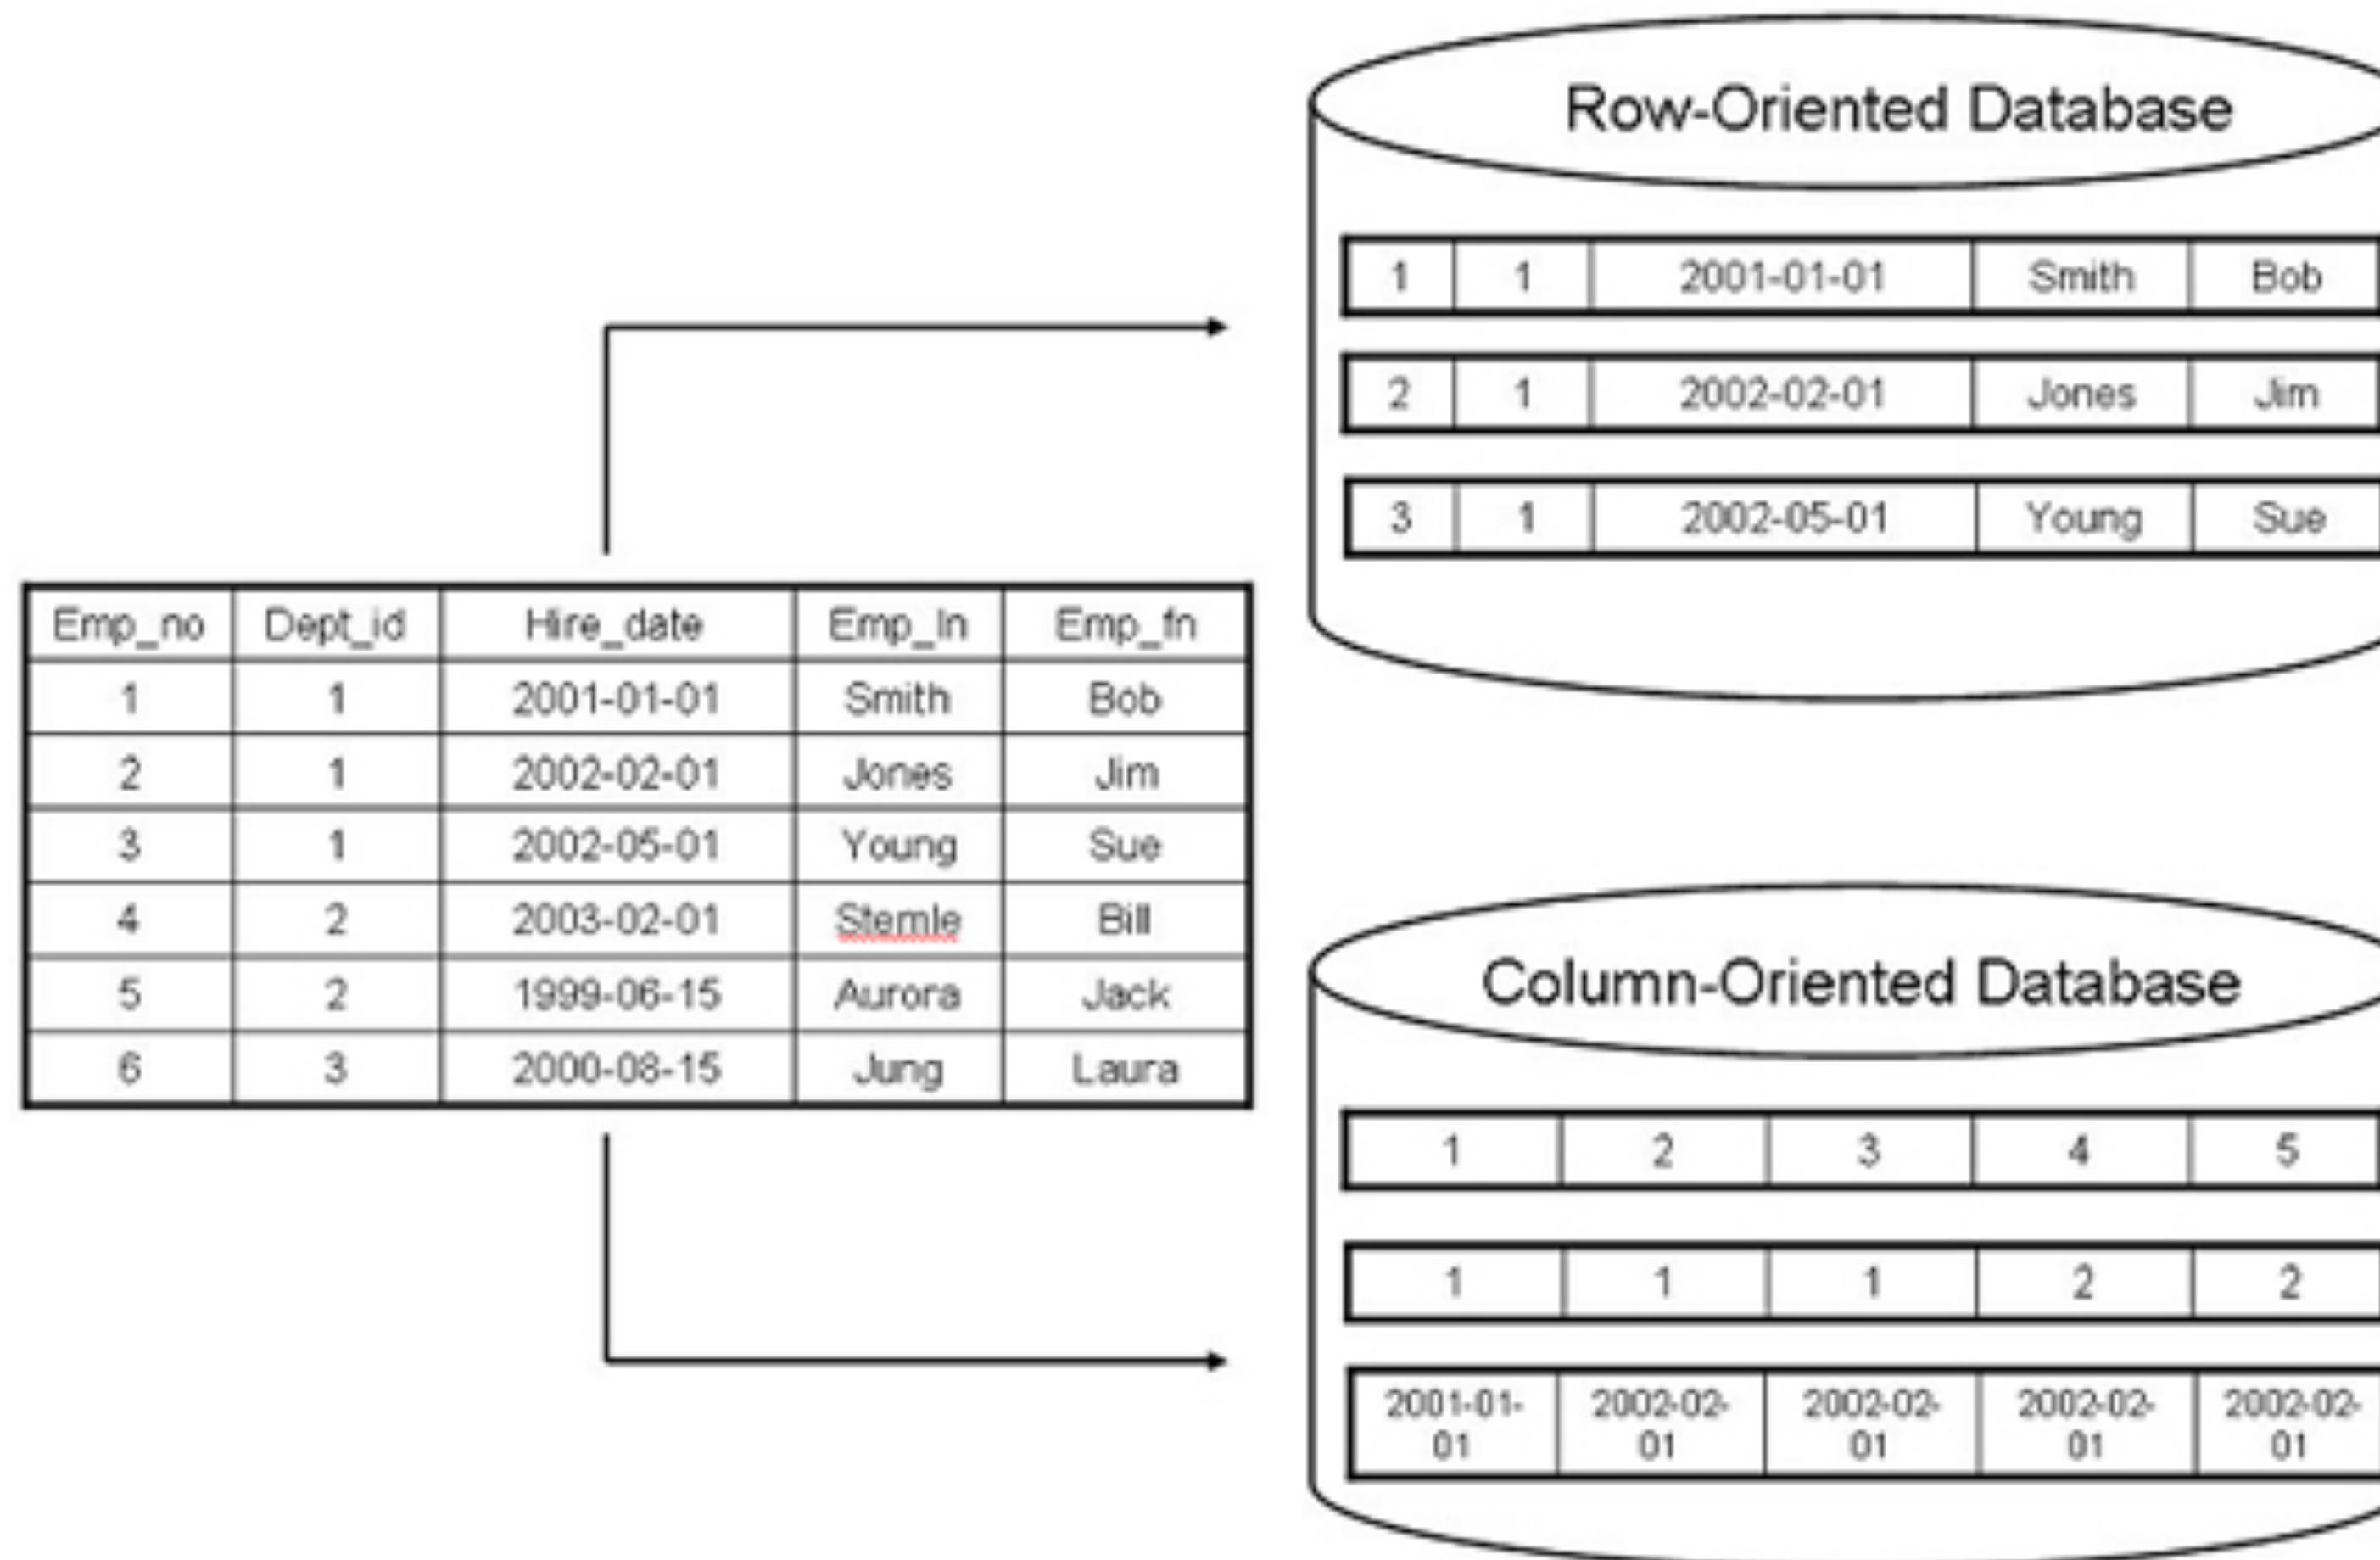
•

4

Features of LevelDB

- Very high speed write: random writes=400,000/s, random reads=60,000/s
- LevelDb is different from Redis as it persists all data in disks.
- LevelDB stores data based on the values of keys. Therefore, close keys are stored together
- LevelDB only supports simple read/write operations.
- LevelDb support snapshots. So you can keep a version of data.

Row Based vs Column Based



Storage of Column Family

Row Key	Students			Branch	
StudentID	Name	Age	Bname	GPA	
100	Ram	18	CSE	7.9	
101	Sham	17	ECE	8	
102	John	18	EEE	7.5	
103	Sam	17	CSE	8.5	

Column

Cells

Row Key

Column Families

KeySpace

Column Family

Key	Column Name	Column Name	Column Name
	Value	Value	Value

Key	Column Name	Column Name
	Value	Value

Key	Column Name	Column Name	Column Name	Column Name
	Value	Value	Value	Value

Sorted by Key

Column

Column Family

Key	Column Name	Column Name	Column Name
	Value	Value	Value
Key	Column Name	Column Name	
	Value	Value	

Sorted by Key

KeySpace

--

- Table with single-row partitions

partition key

columns

partitions

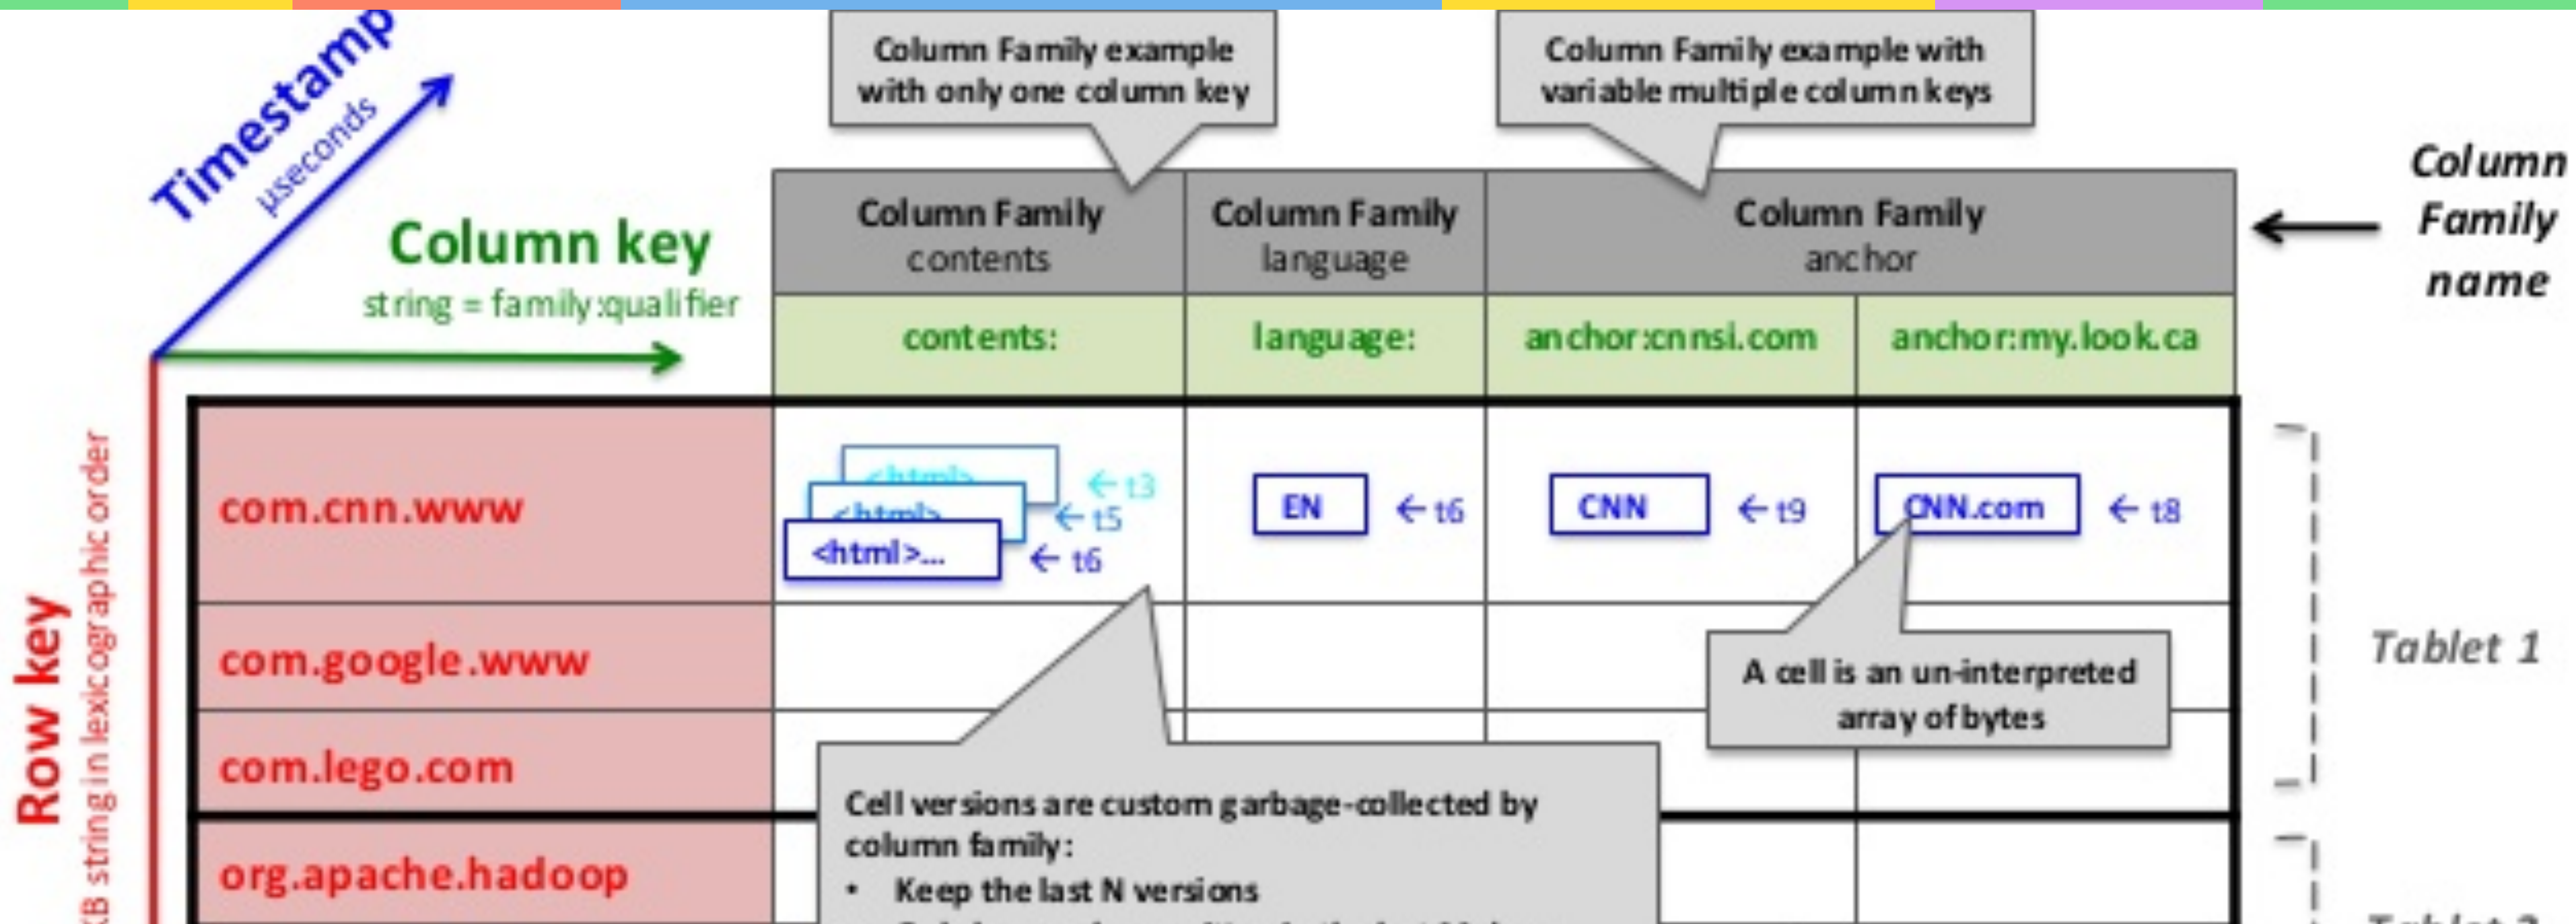
rows

cells

performer	born	country	died	founded	style	type
John Lennon	1940	England	1980		Rock	artist
Paul McCartney	1942	England			Rock	artist
The Beatles		England		1957	Rock	band

- Column family view

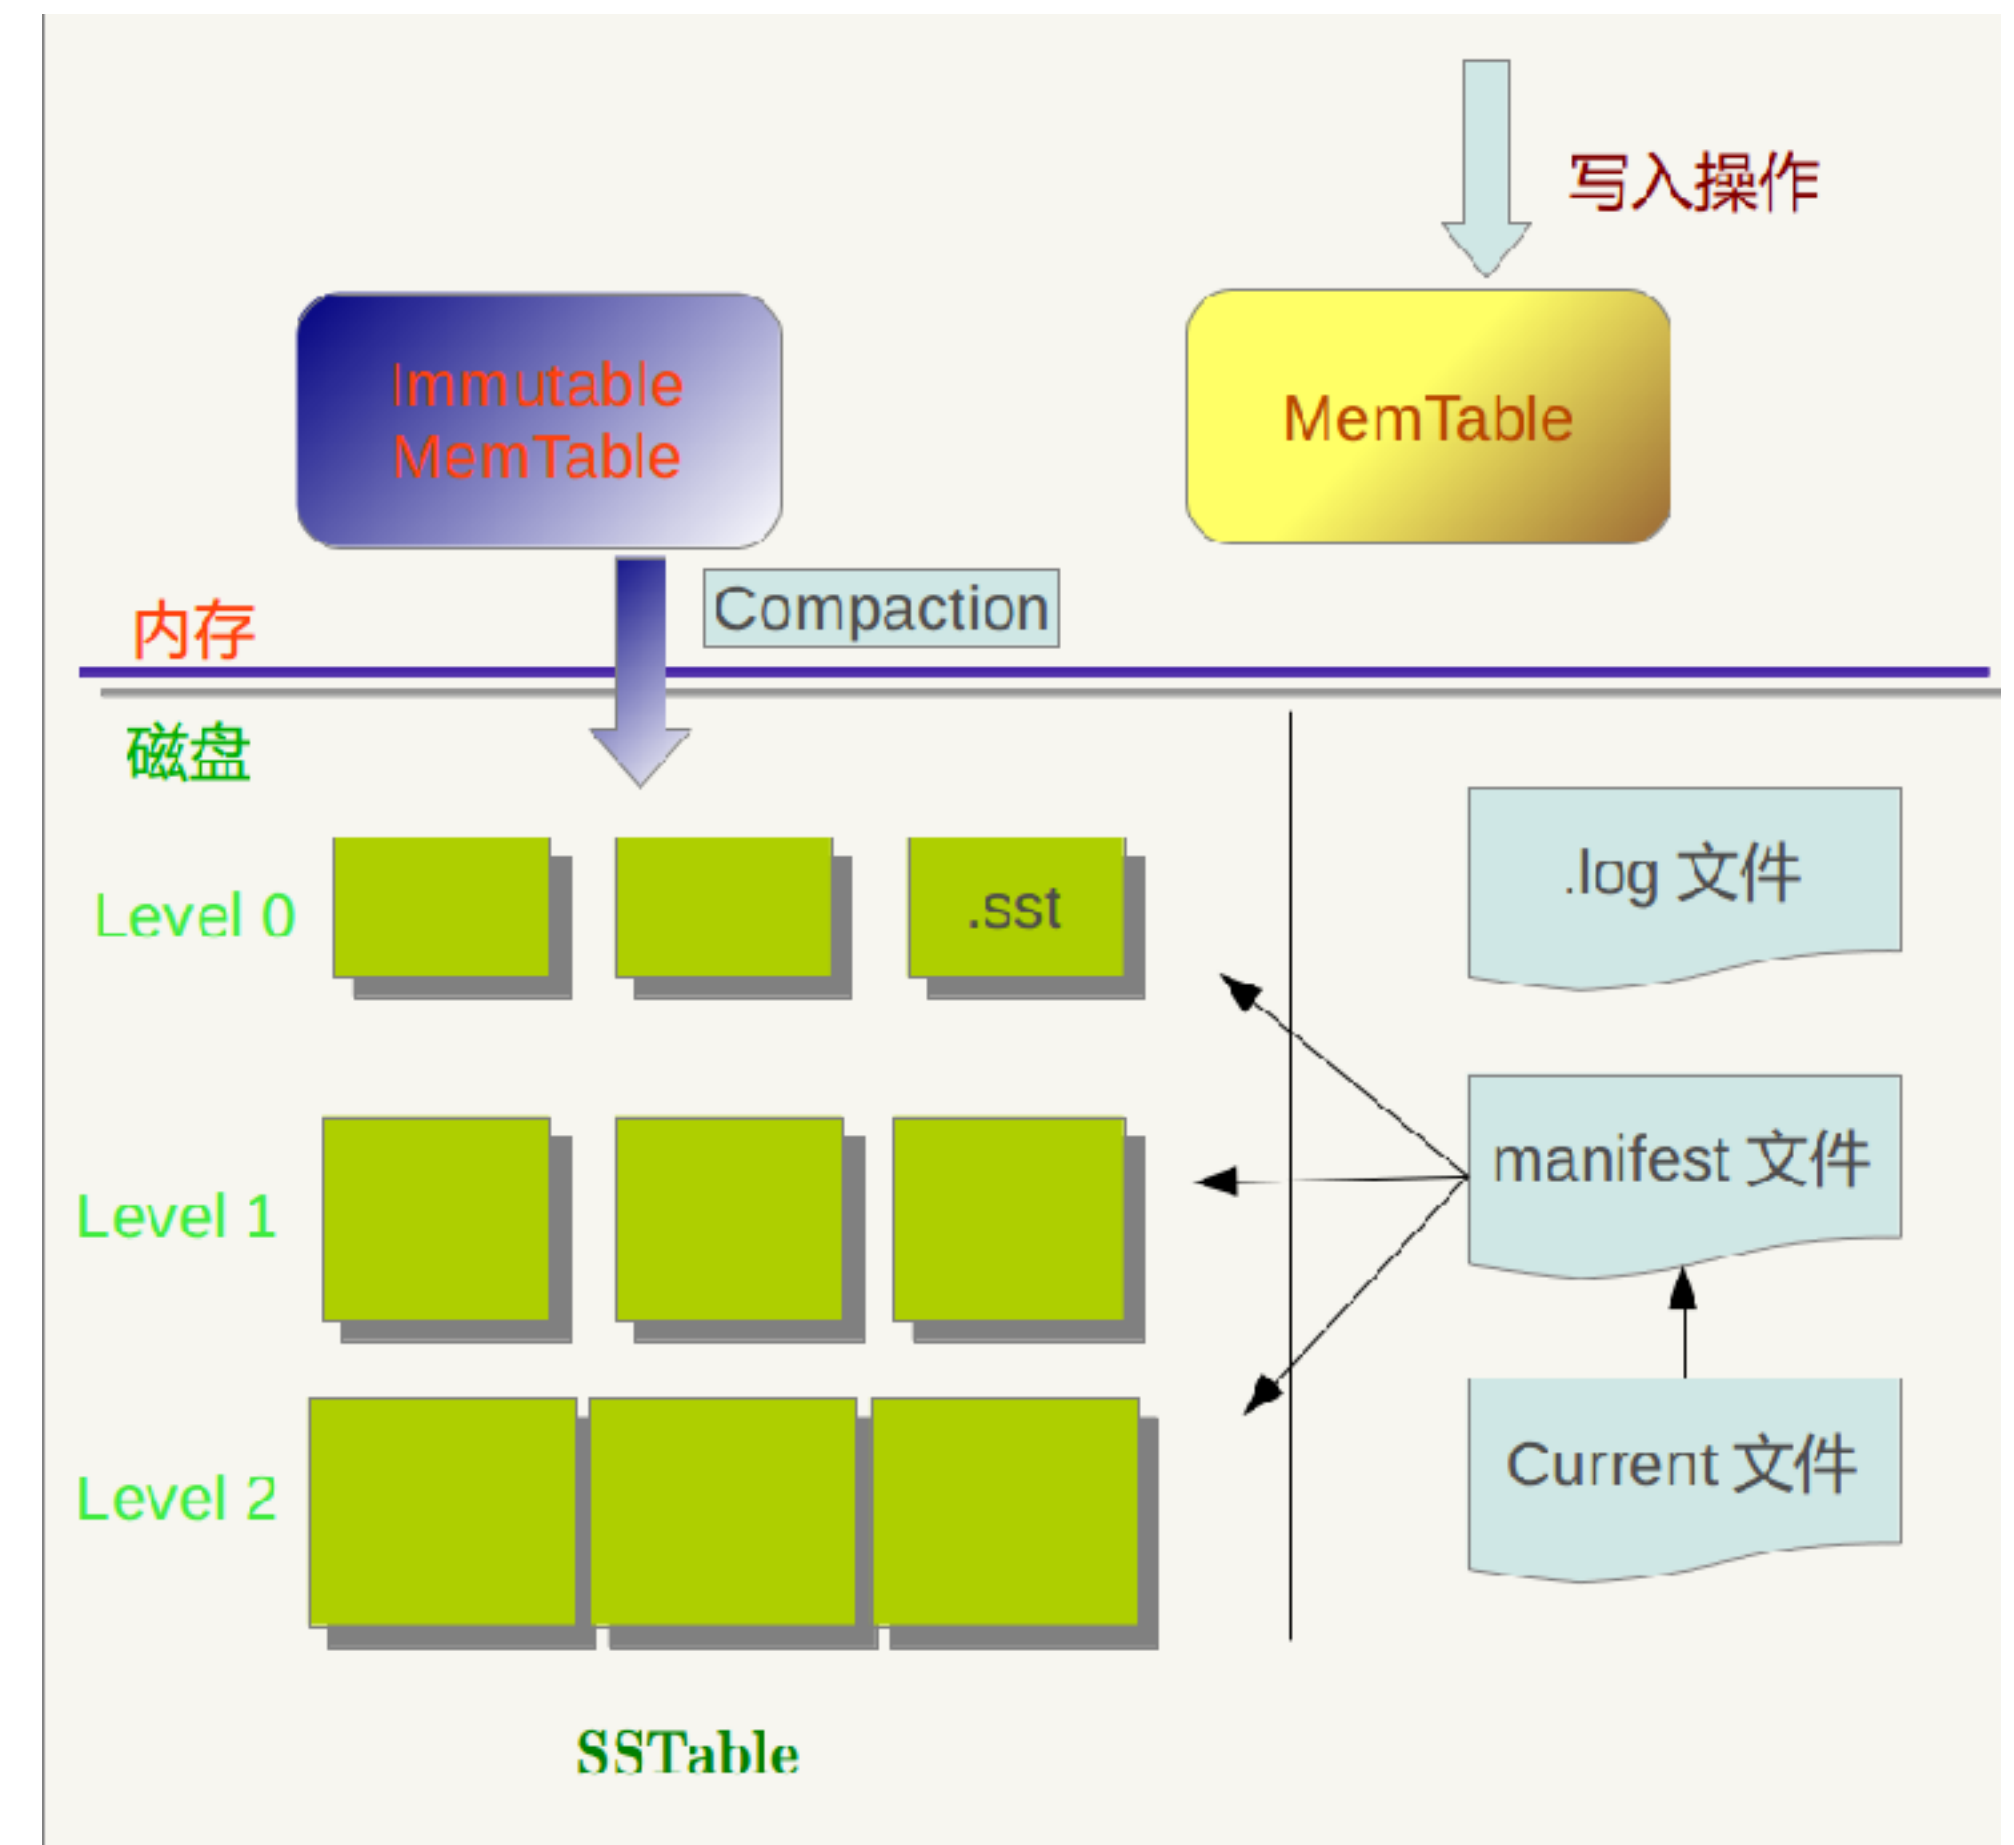
John Lennon	<div>born</div> <div>1940</div>	<div>country</div> <div>England</div>	<div>died</div> <div>1980</div>		<div>style</div> <div>Rock</div>	<div>type</div> <div>artist</div>
Paul McCartney	<div>born</div> <div>1942</div>	<div>country</div> <div>England</div>			<div>style</div> <div>Rock</div>	<div>type</div> <div>artist</div>
The Beatles		<div>country</div> <div>England</div>		<div>founded</div> <div>1957</div>	<div>style</div> <div>Rock</div>	<div>type</div> <div>band</div>



The Idea of Key-Value Format

Architecture of LevelDB

- LevelDB has many layers
- In memory, there is an LSM (Log-Structured Merge tree) to maintain index. If the LSM is full, it will be locked as a Memtable.
- Disk part is partitioned into Level 0 - Level K. Each level is N (N=10 by default) times larger than previous one.
- If one level is full, it needs to be flushed and merged with the next level. The process is called compaction



API of LevelDB

```
DB() { };  
virtual ~DB();  
static Status Open(const Options& options,  
                  const std::string& name,  
                  DB** dbptr);  
virtual Status Put(const WriteOptions& options,  
                  const Slice& key,  
                  const Slice& value) = 0;  
virtual Status Delete(const WriteOptions& options, const  
Slice& key) = 0;  
virtual Status Write(const WriteOptions& options, WriteBatch*  
updates) = 0;
```

```
virtual Status Get(const ReadOptions& options,  
                  const Slice& key, std::string* value) = 0;  
virtual Iterator* NewIterator(const ReadOptions& options) = 0;  
virtual const Snapshot* GetSnapshot() = 0;  
virtual void ReleaseSnapshot(const Snapshot* snapshot) = 0;
```


In Memory Table: MemTable


- All read/writes are first buffered in the MemTable
- By default, the size of a MemTable is 4M
- When a MemTable is full, the MemTable is locked. New updates will be sent to a newly created MemTable. The old one waits for compaction.

one question: Given the MemTable, how to efficiently locate the data?

We need an Index!

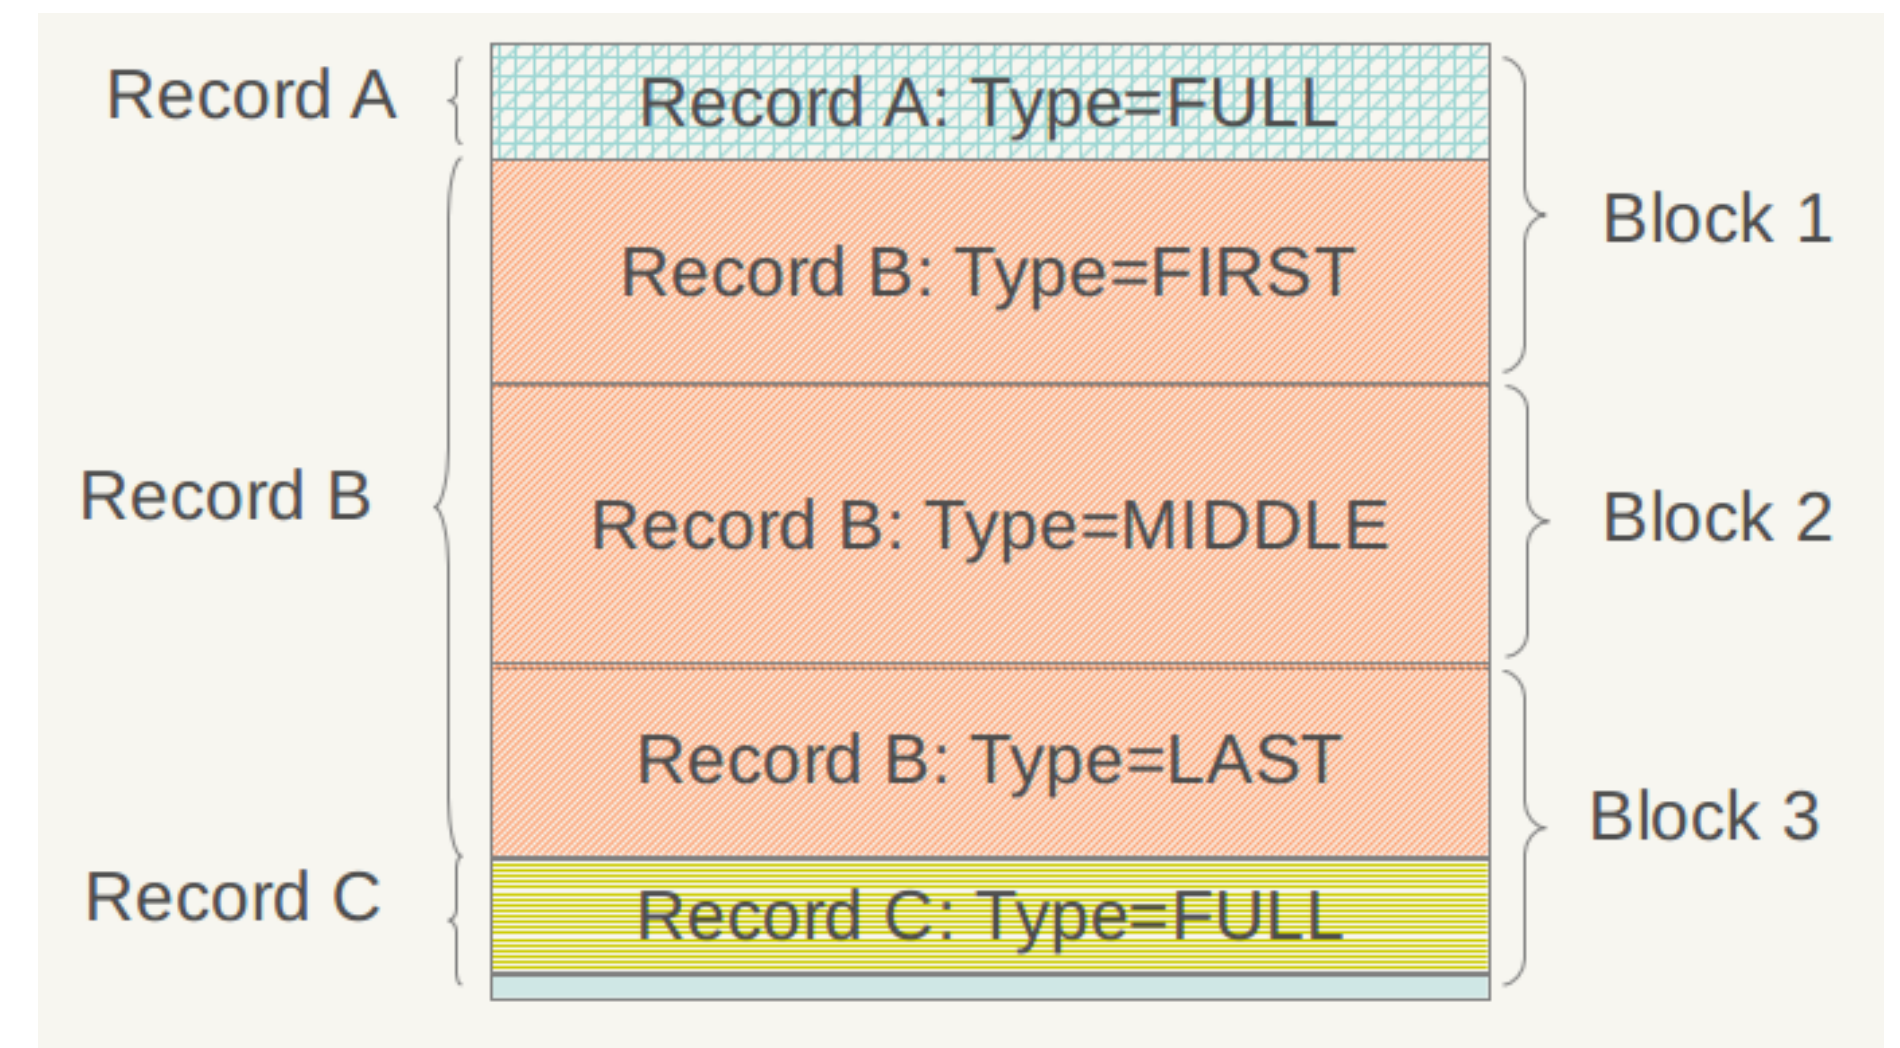

MemTable=SkipList

- Similar to the Binary Tree
- Keep balanced with a high probability
- No need rotations in the tree to keep balanced




Do not Forget the Log

- To avoid data loss, before update data in memory, we write create a log record.
- In LevelDB, the format of log is:
- CRC32 is the verification code, length= record length, type=full, first, middle, last



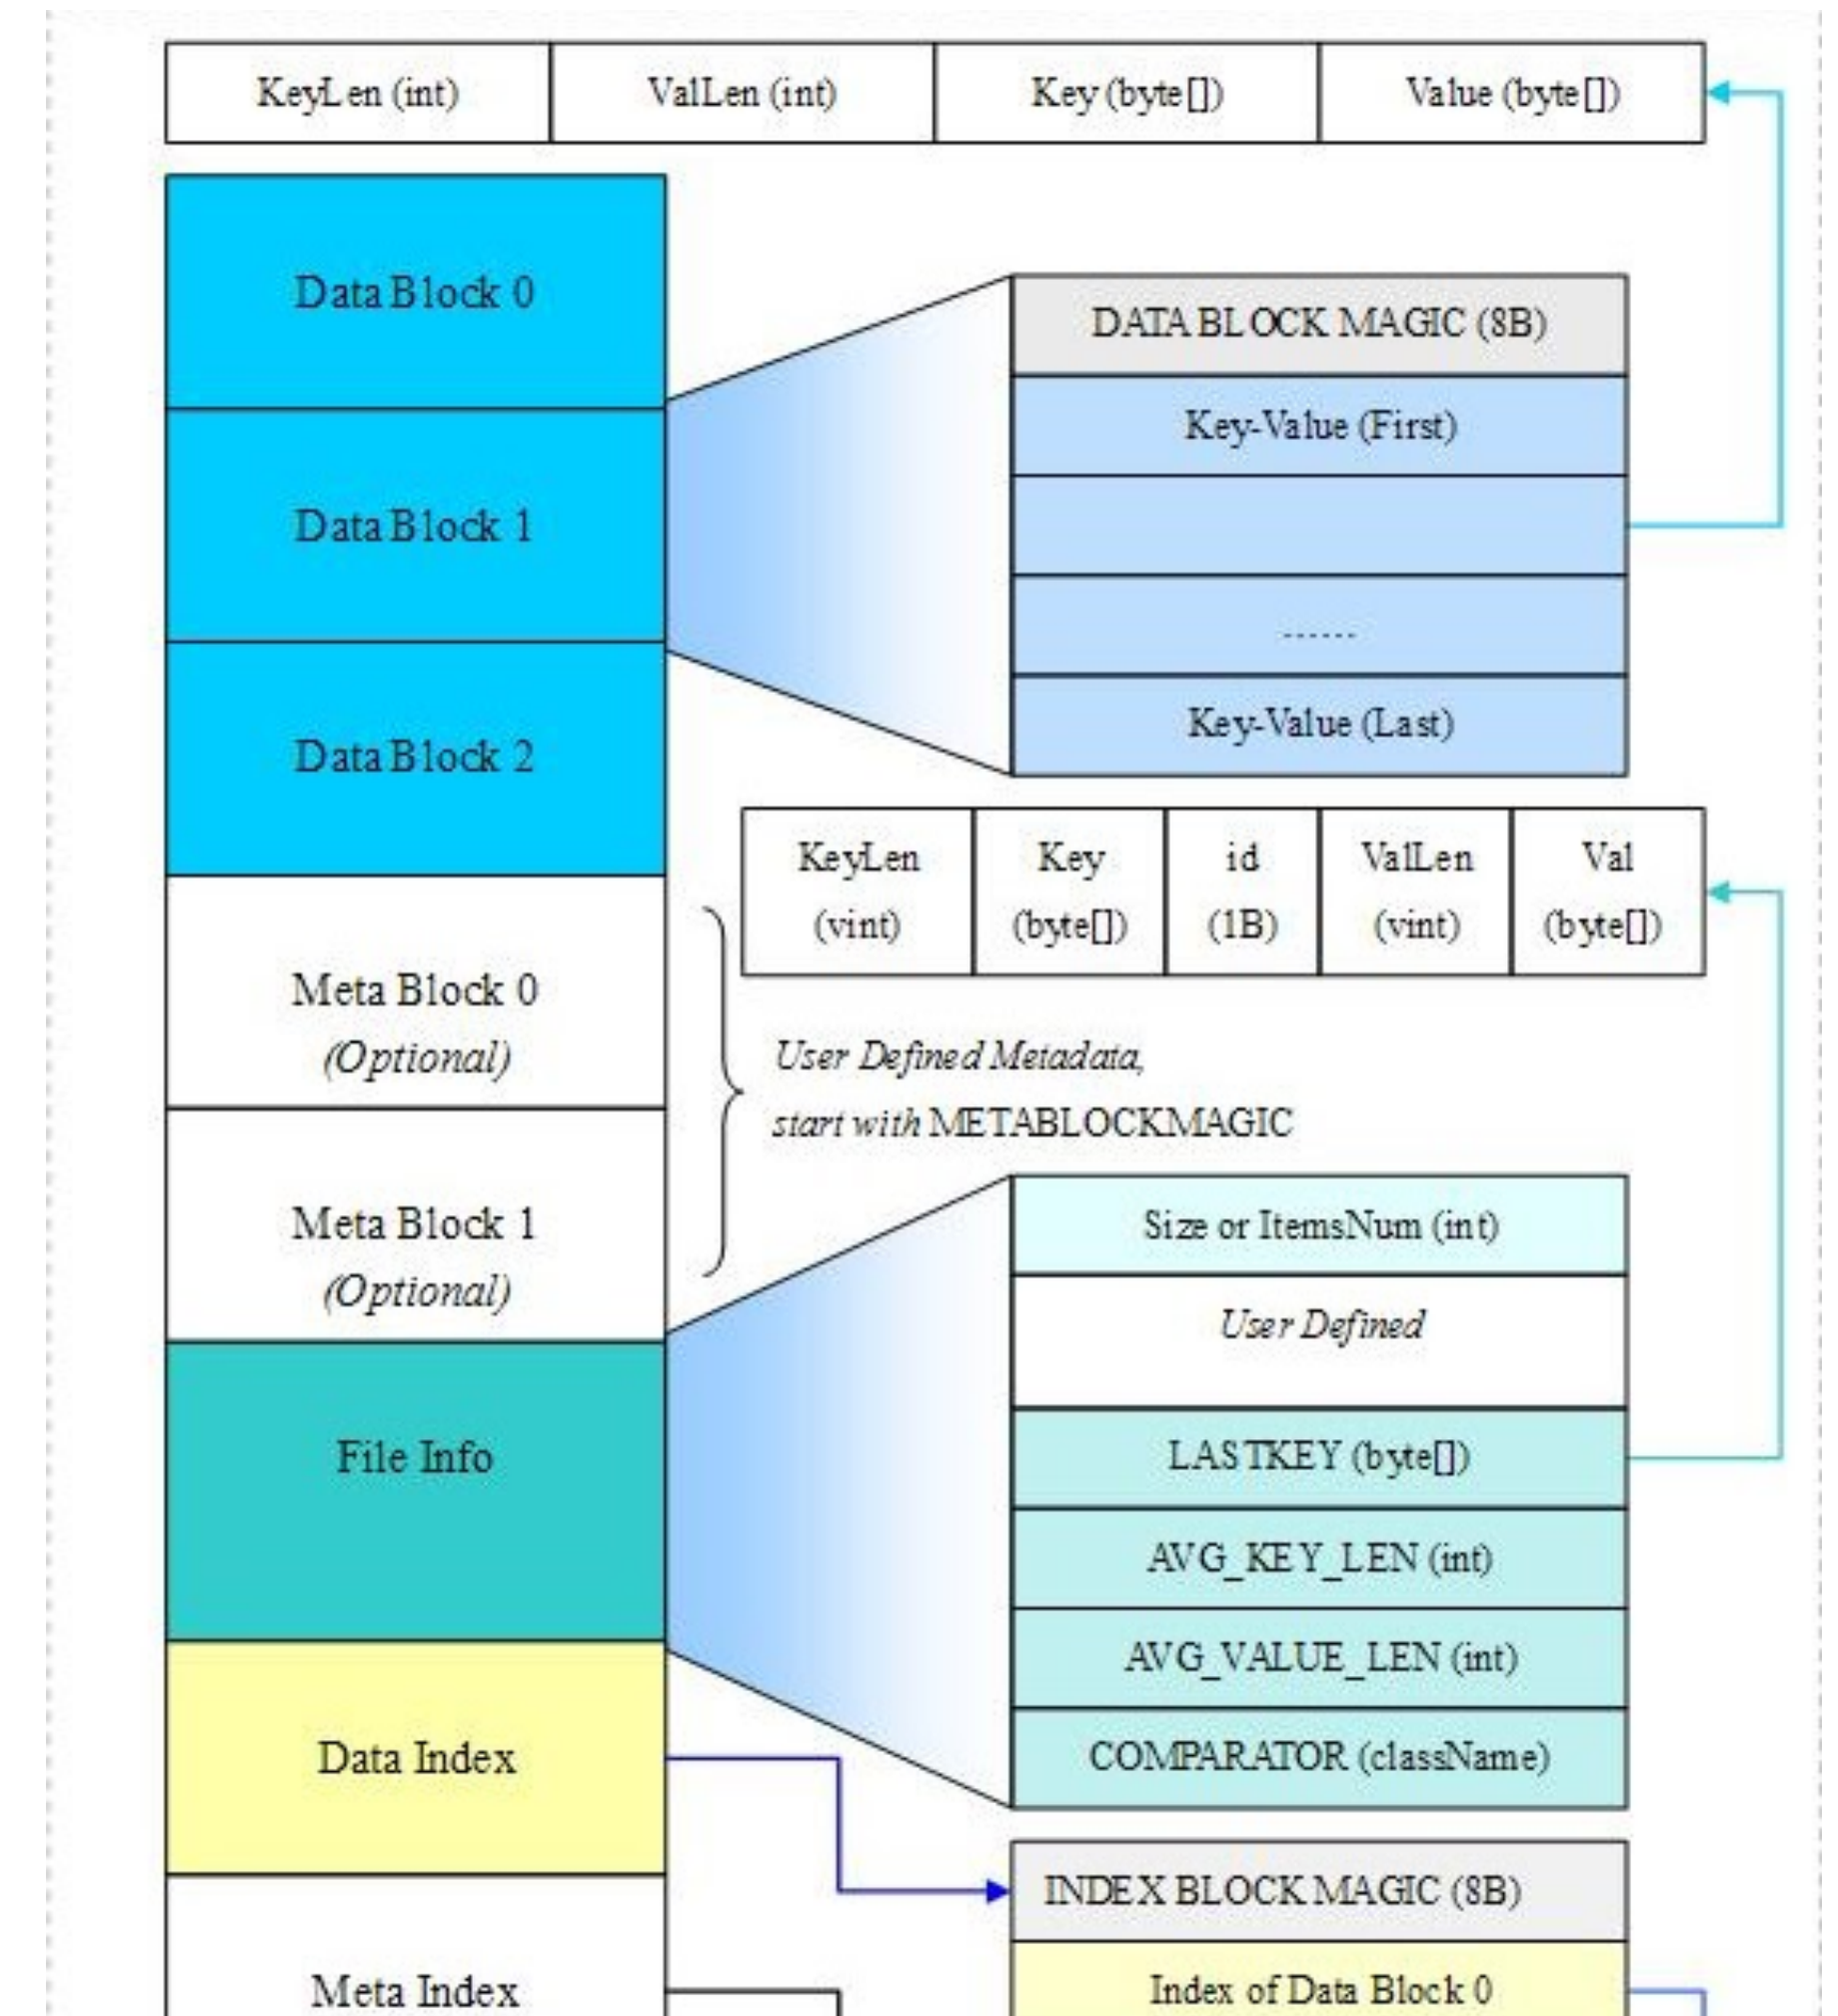
MemTable Updates

- If MemTable is full, we need to flush it to the disk.
- LevelDb will create a new Memtable and the corresponding log file. The old one is called the Immutable Memtable, which will be written as a SSTable in the disk.

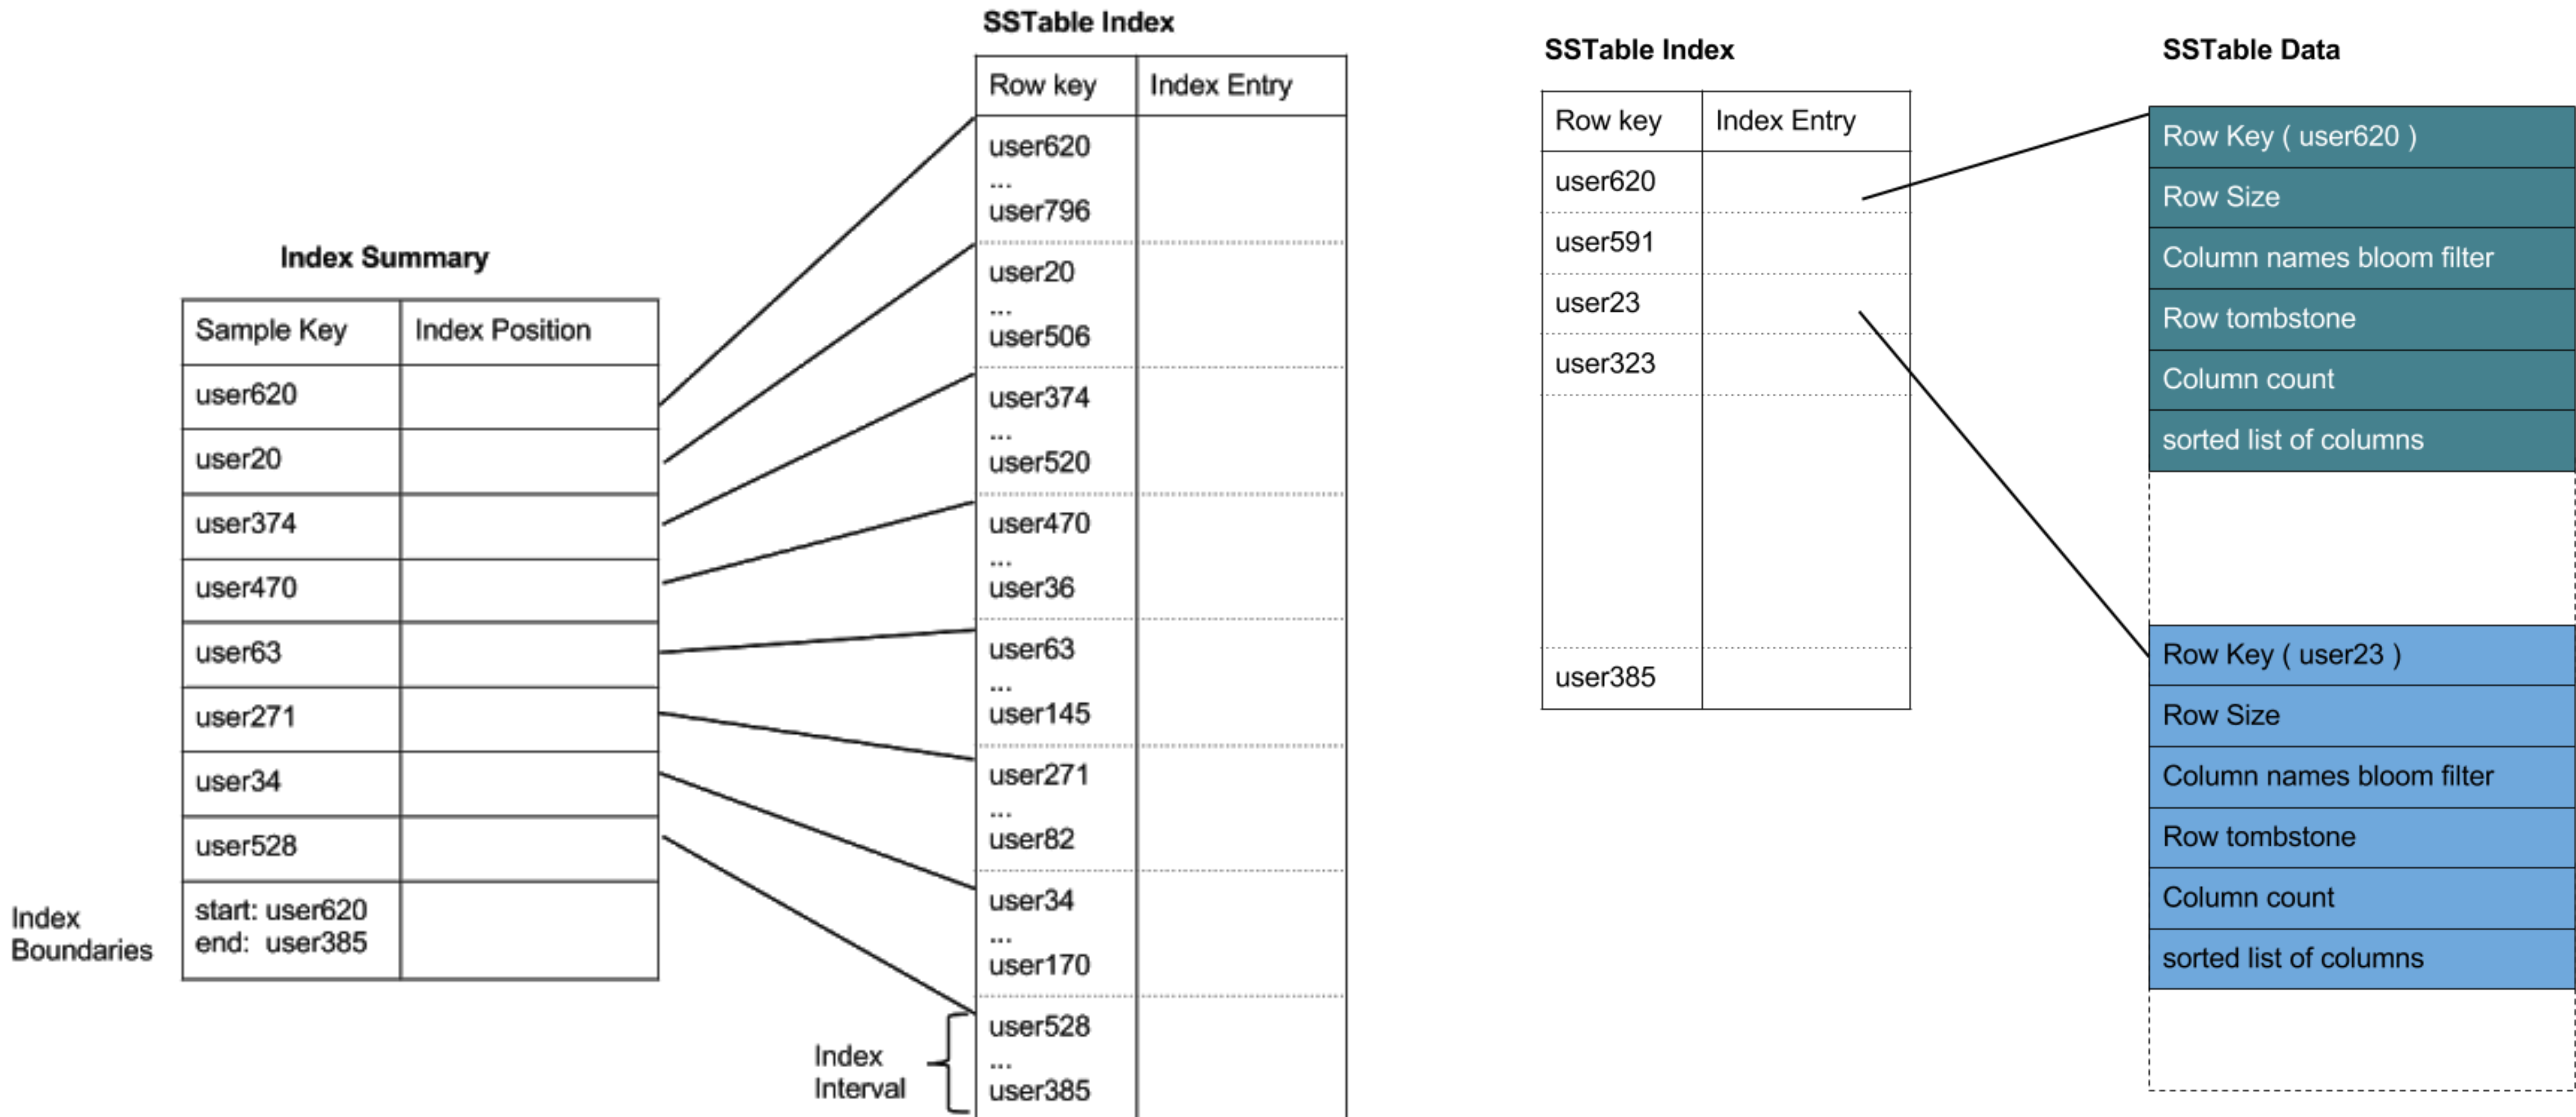


Disk File: SSTable

- SSTable is the disk part of LevelDB. Each sstable is 2MB and belongs to one specific level:
- level 0: at most 4 sstables
- level 1: the total size of sstable less than 10M
- level 2: the total size of sstable less than 1000M

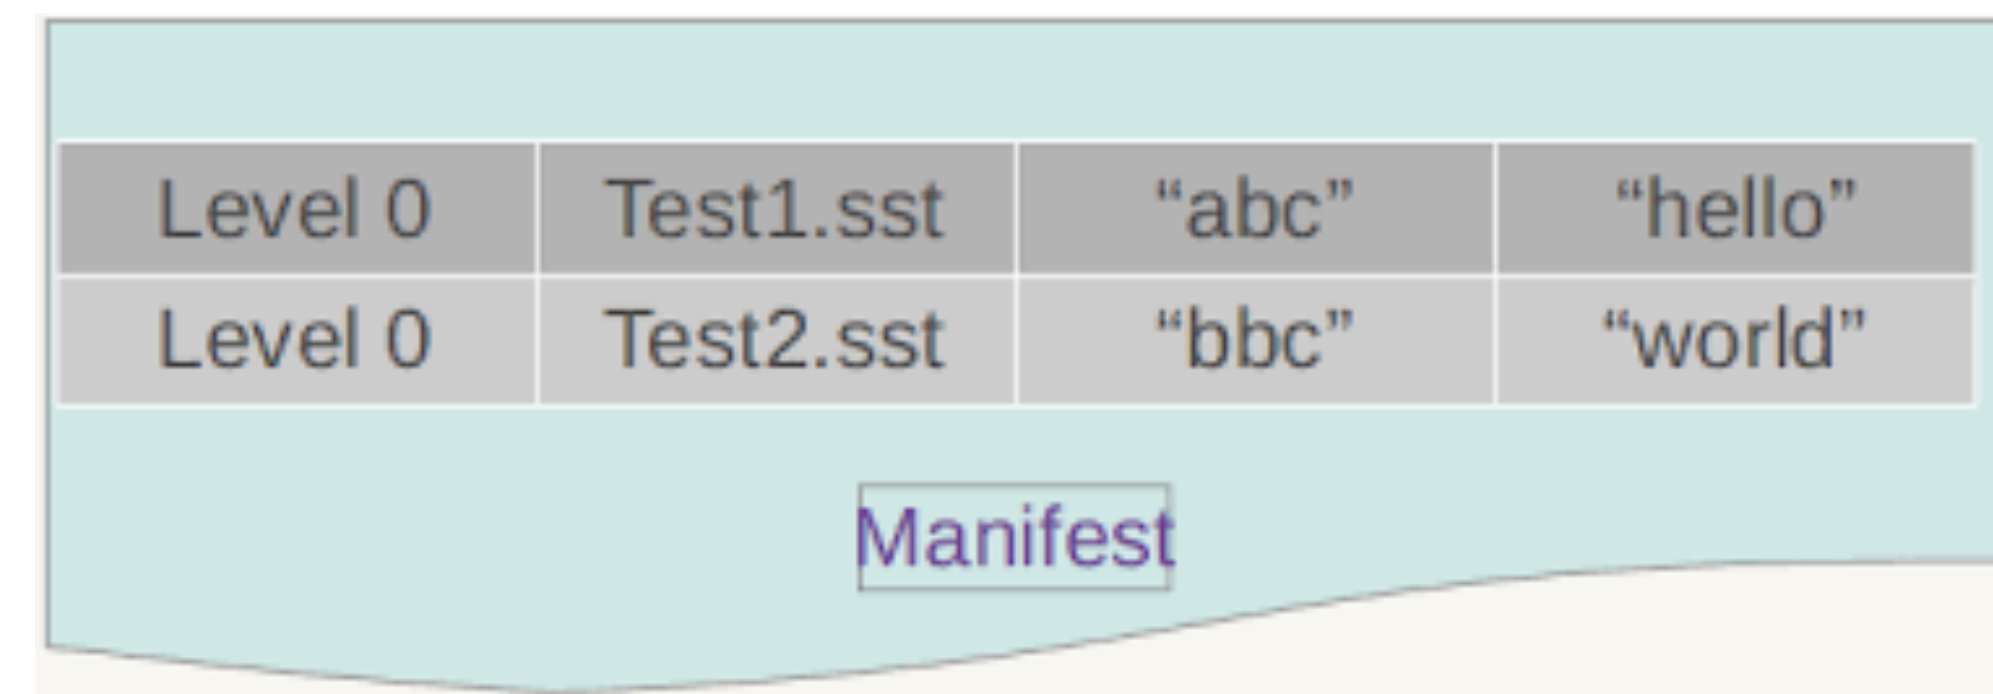


BlockIndex



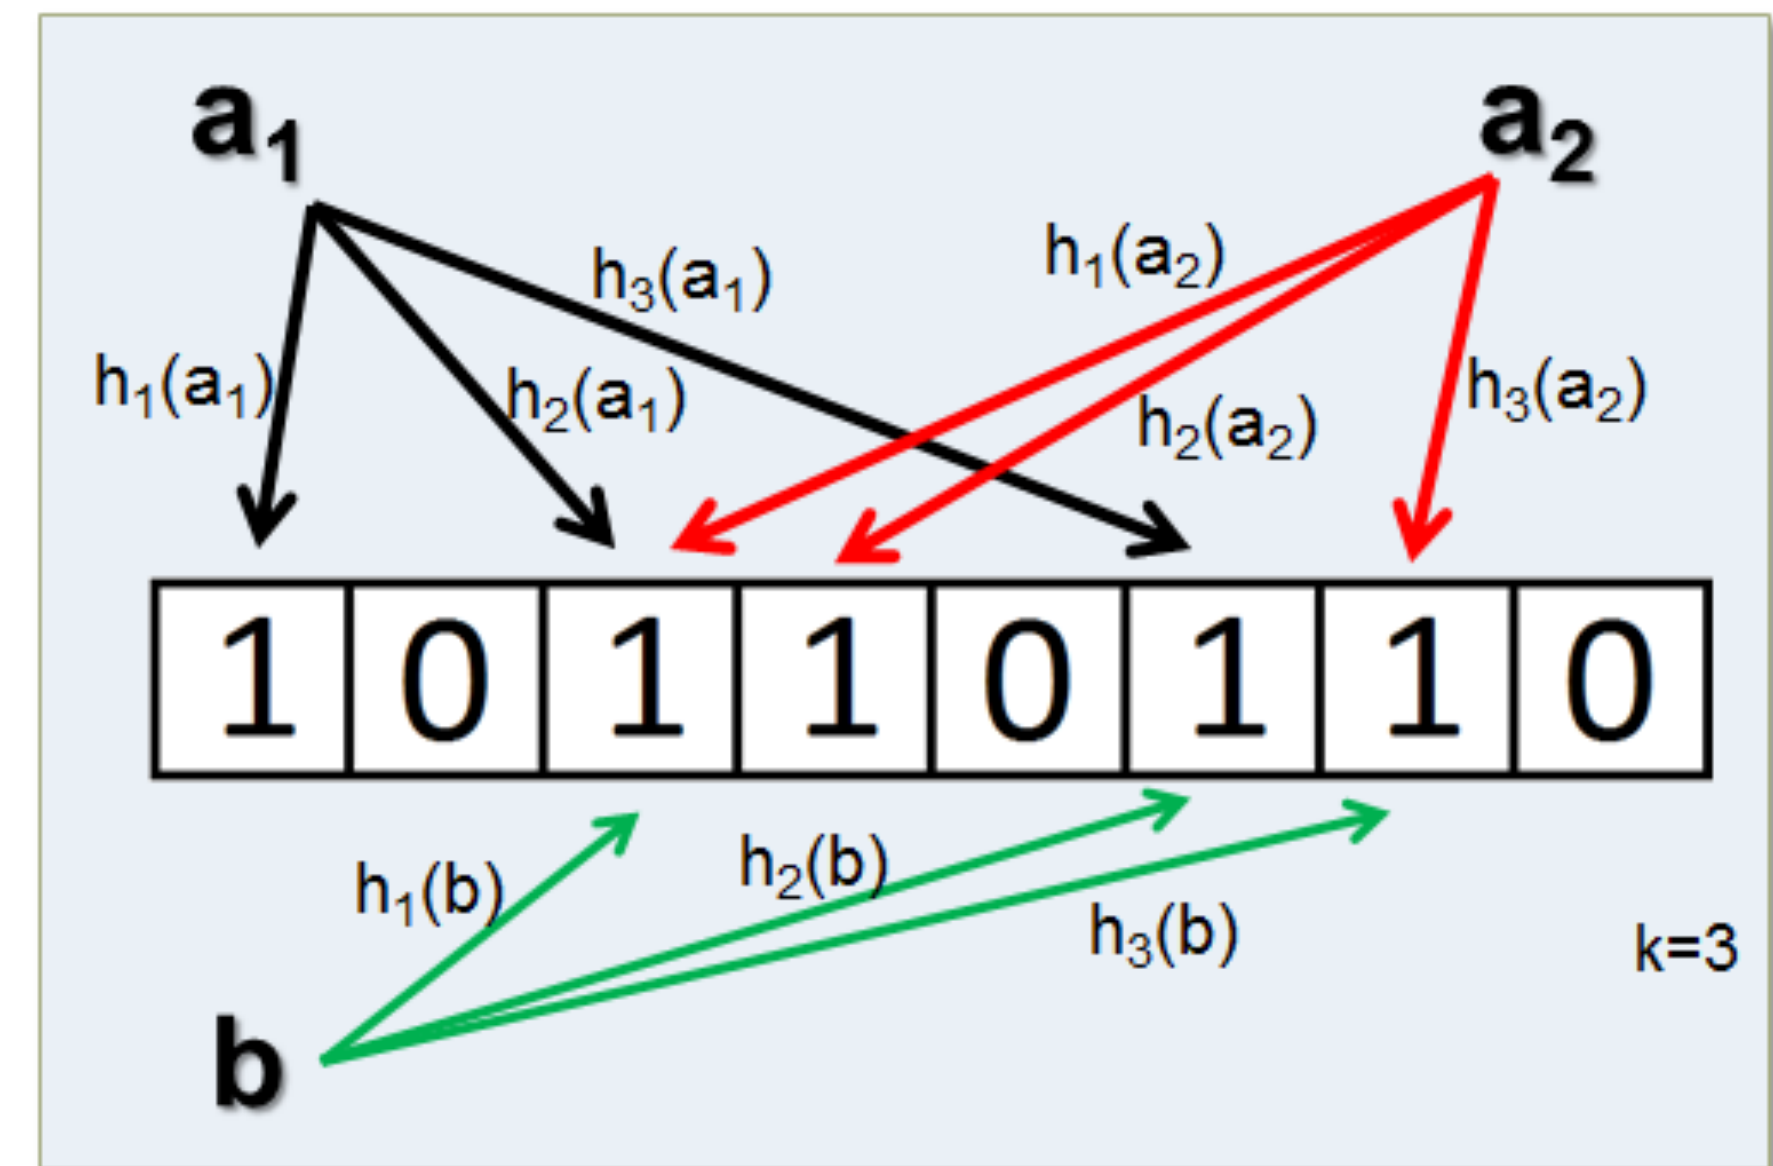
Meta-Index

- The search problem remains. We cannot afford searching every SSTable for a single key.
- We need to know whether a key DOES exist in a block or not!
- One specific index: BloomFilter is introduced



Bloomfilter

- Bloomfilter=bitmap, Initialized all 0s.
suppose we have total m bits.
- We need K independent Hash function. For every value v , we have $h_1(v), h_2(v), \dots, h_k(v)$ hash values, which mapped v into k positions from 0 to m .
- We set values in those K positions to 1



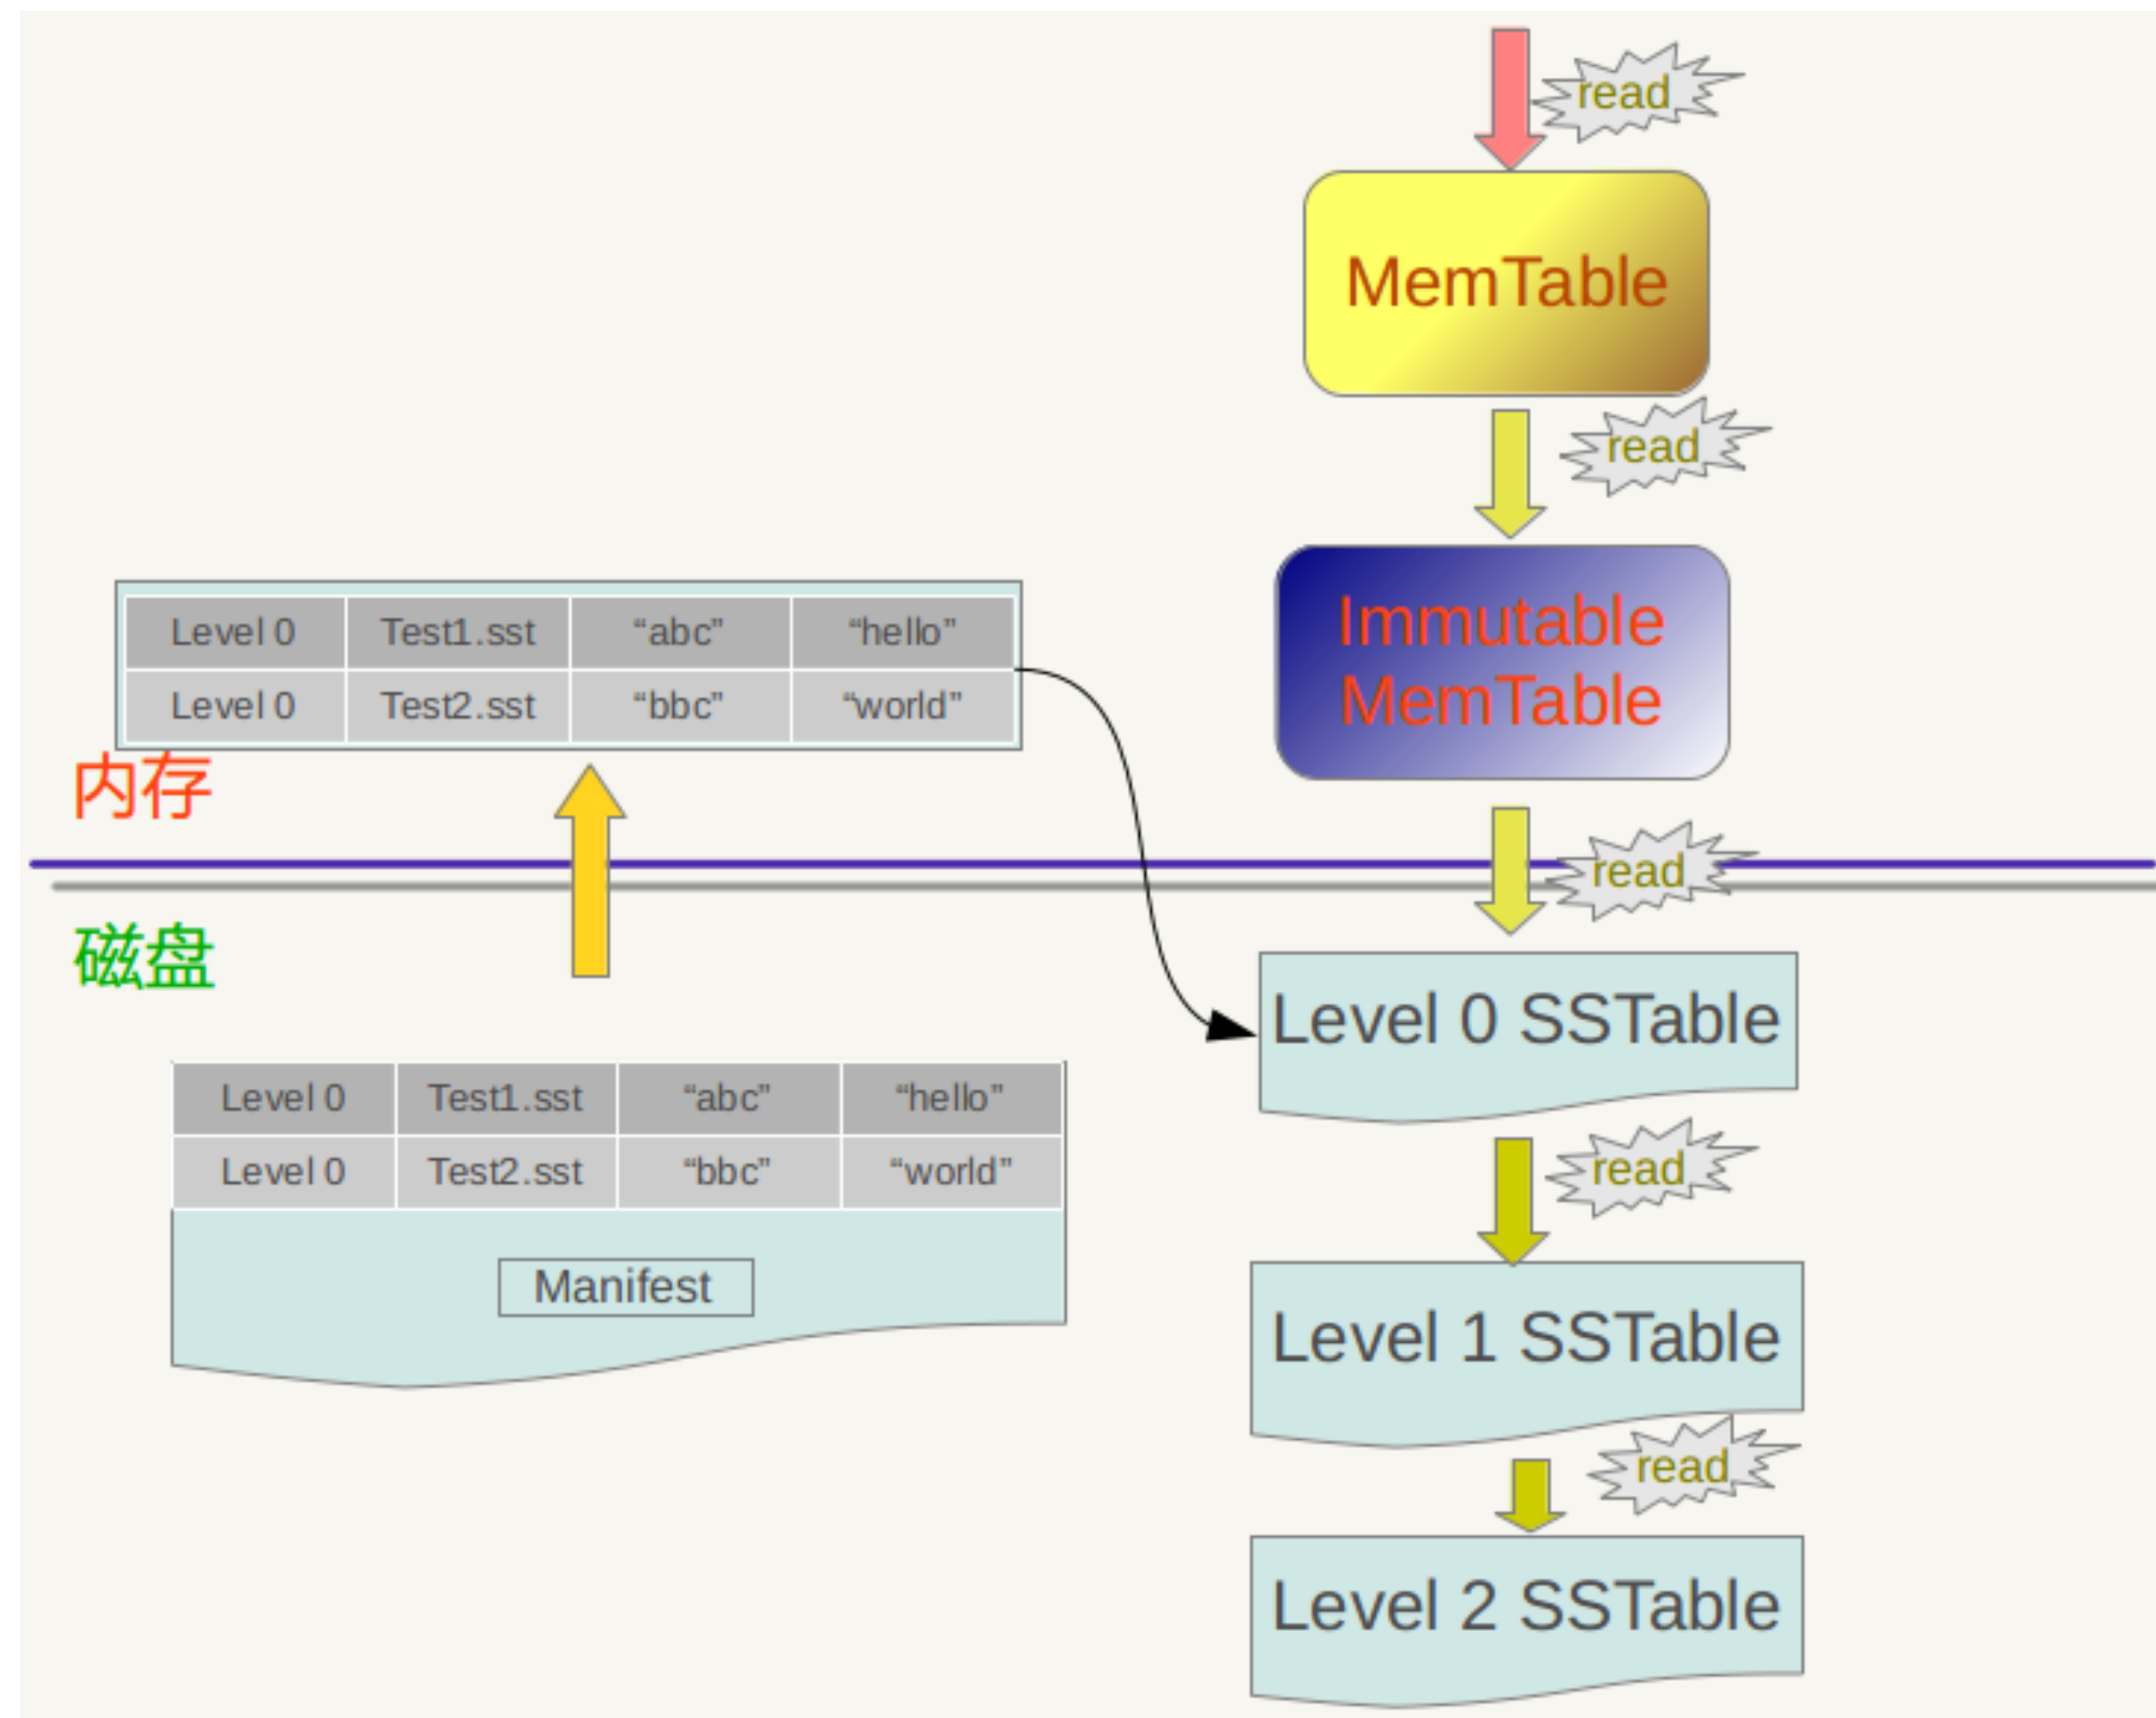
Some Fact of Bloomfilter

- The probability that a certain bit is not set to 1 by a certain hash function during the insertion of an element: $1 - 1/m$
- k is the number of hash functions, the probability that the bit is not set to 1 by any of the hash functions is: $(1 - 1/m)^k$

$$\left(1 - \left[1 - \frac{1}{m}\right]^{kn}\right)^k \approx \left(1 - e^{-kn/m}\right)^k.$$

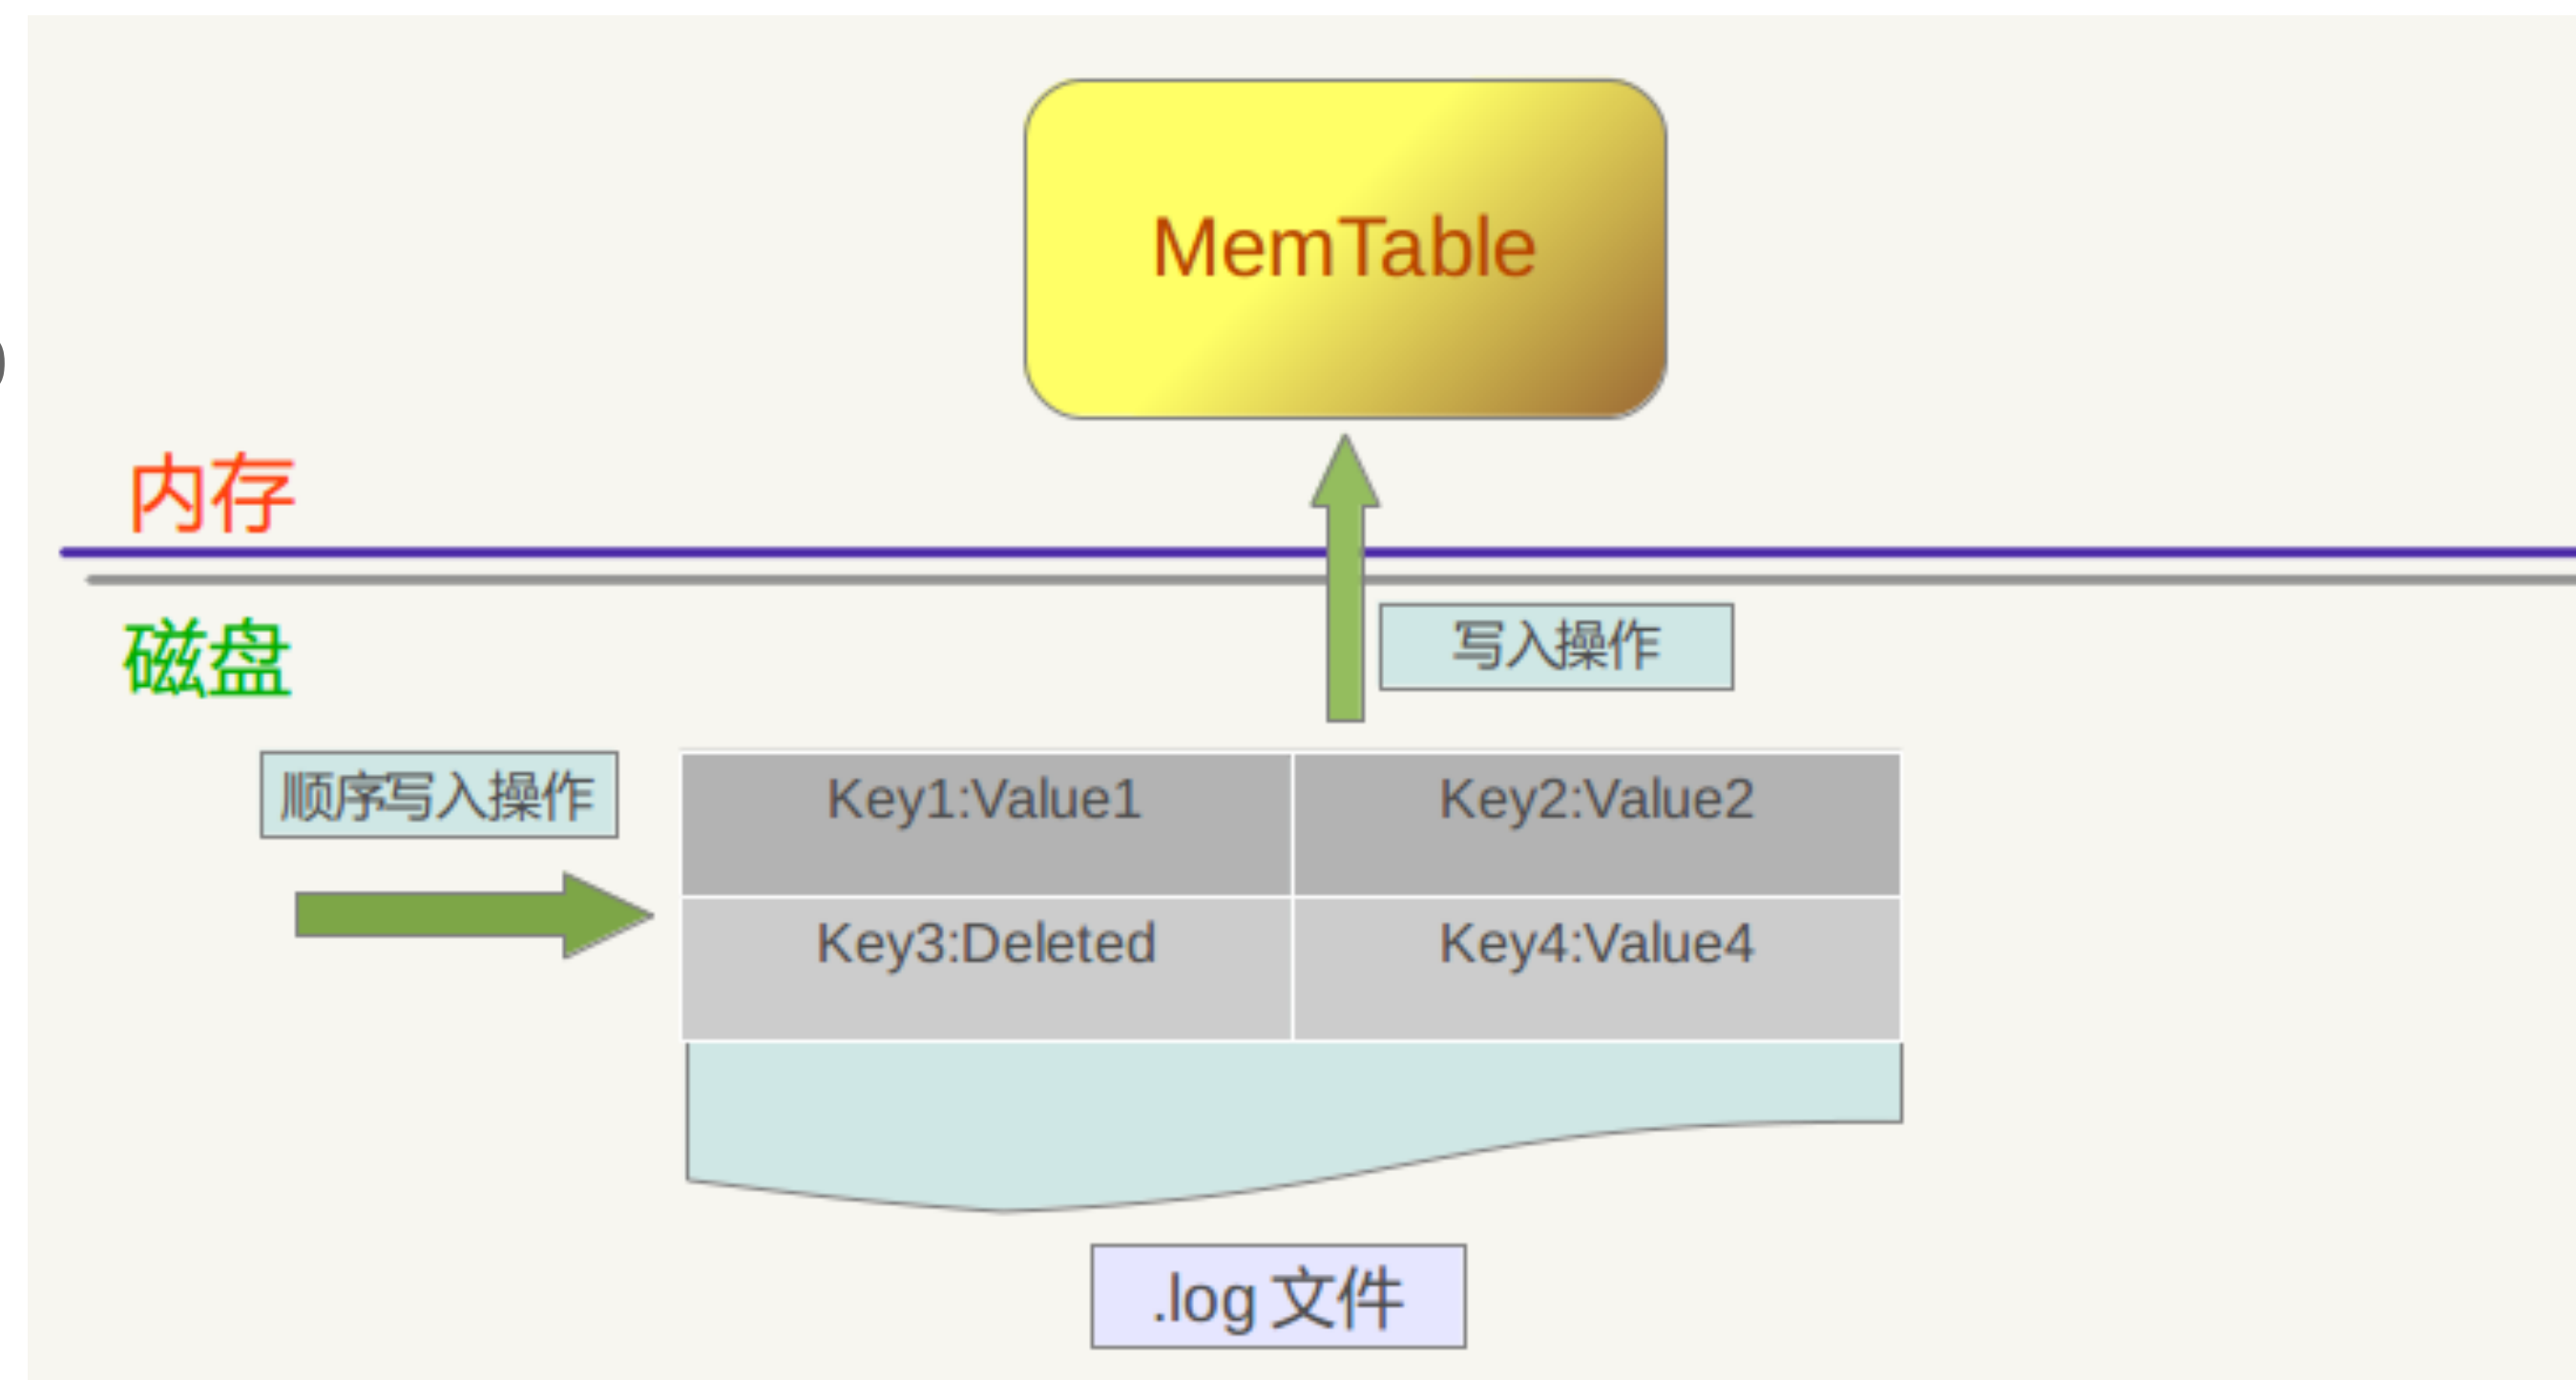
- If we have inserted n elements, the probability that a certain bit is still 0 : $(1 - 1/m)^{nk}$
- the probability that it is 1 is therefore: $1 - (1 - 1/m)^{nk}$
- The probability of all of them being 1, which would cause the algorithm to erroneously claim that the element is in the set, is often given as

Read in LevelDB

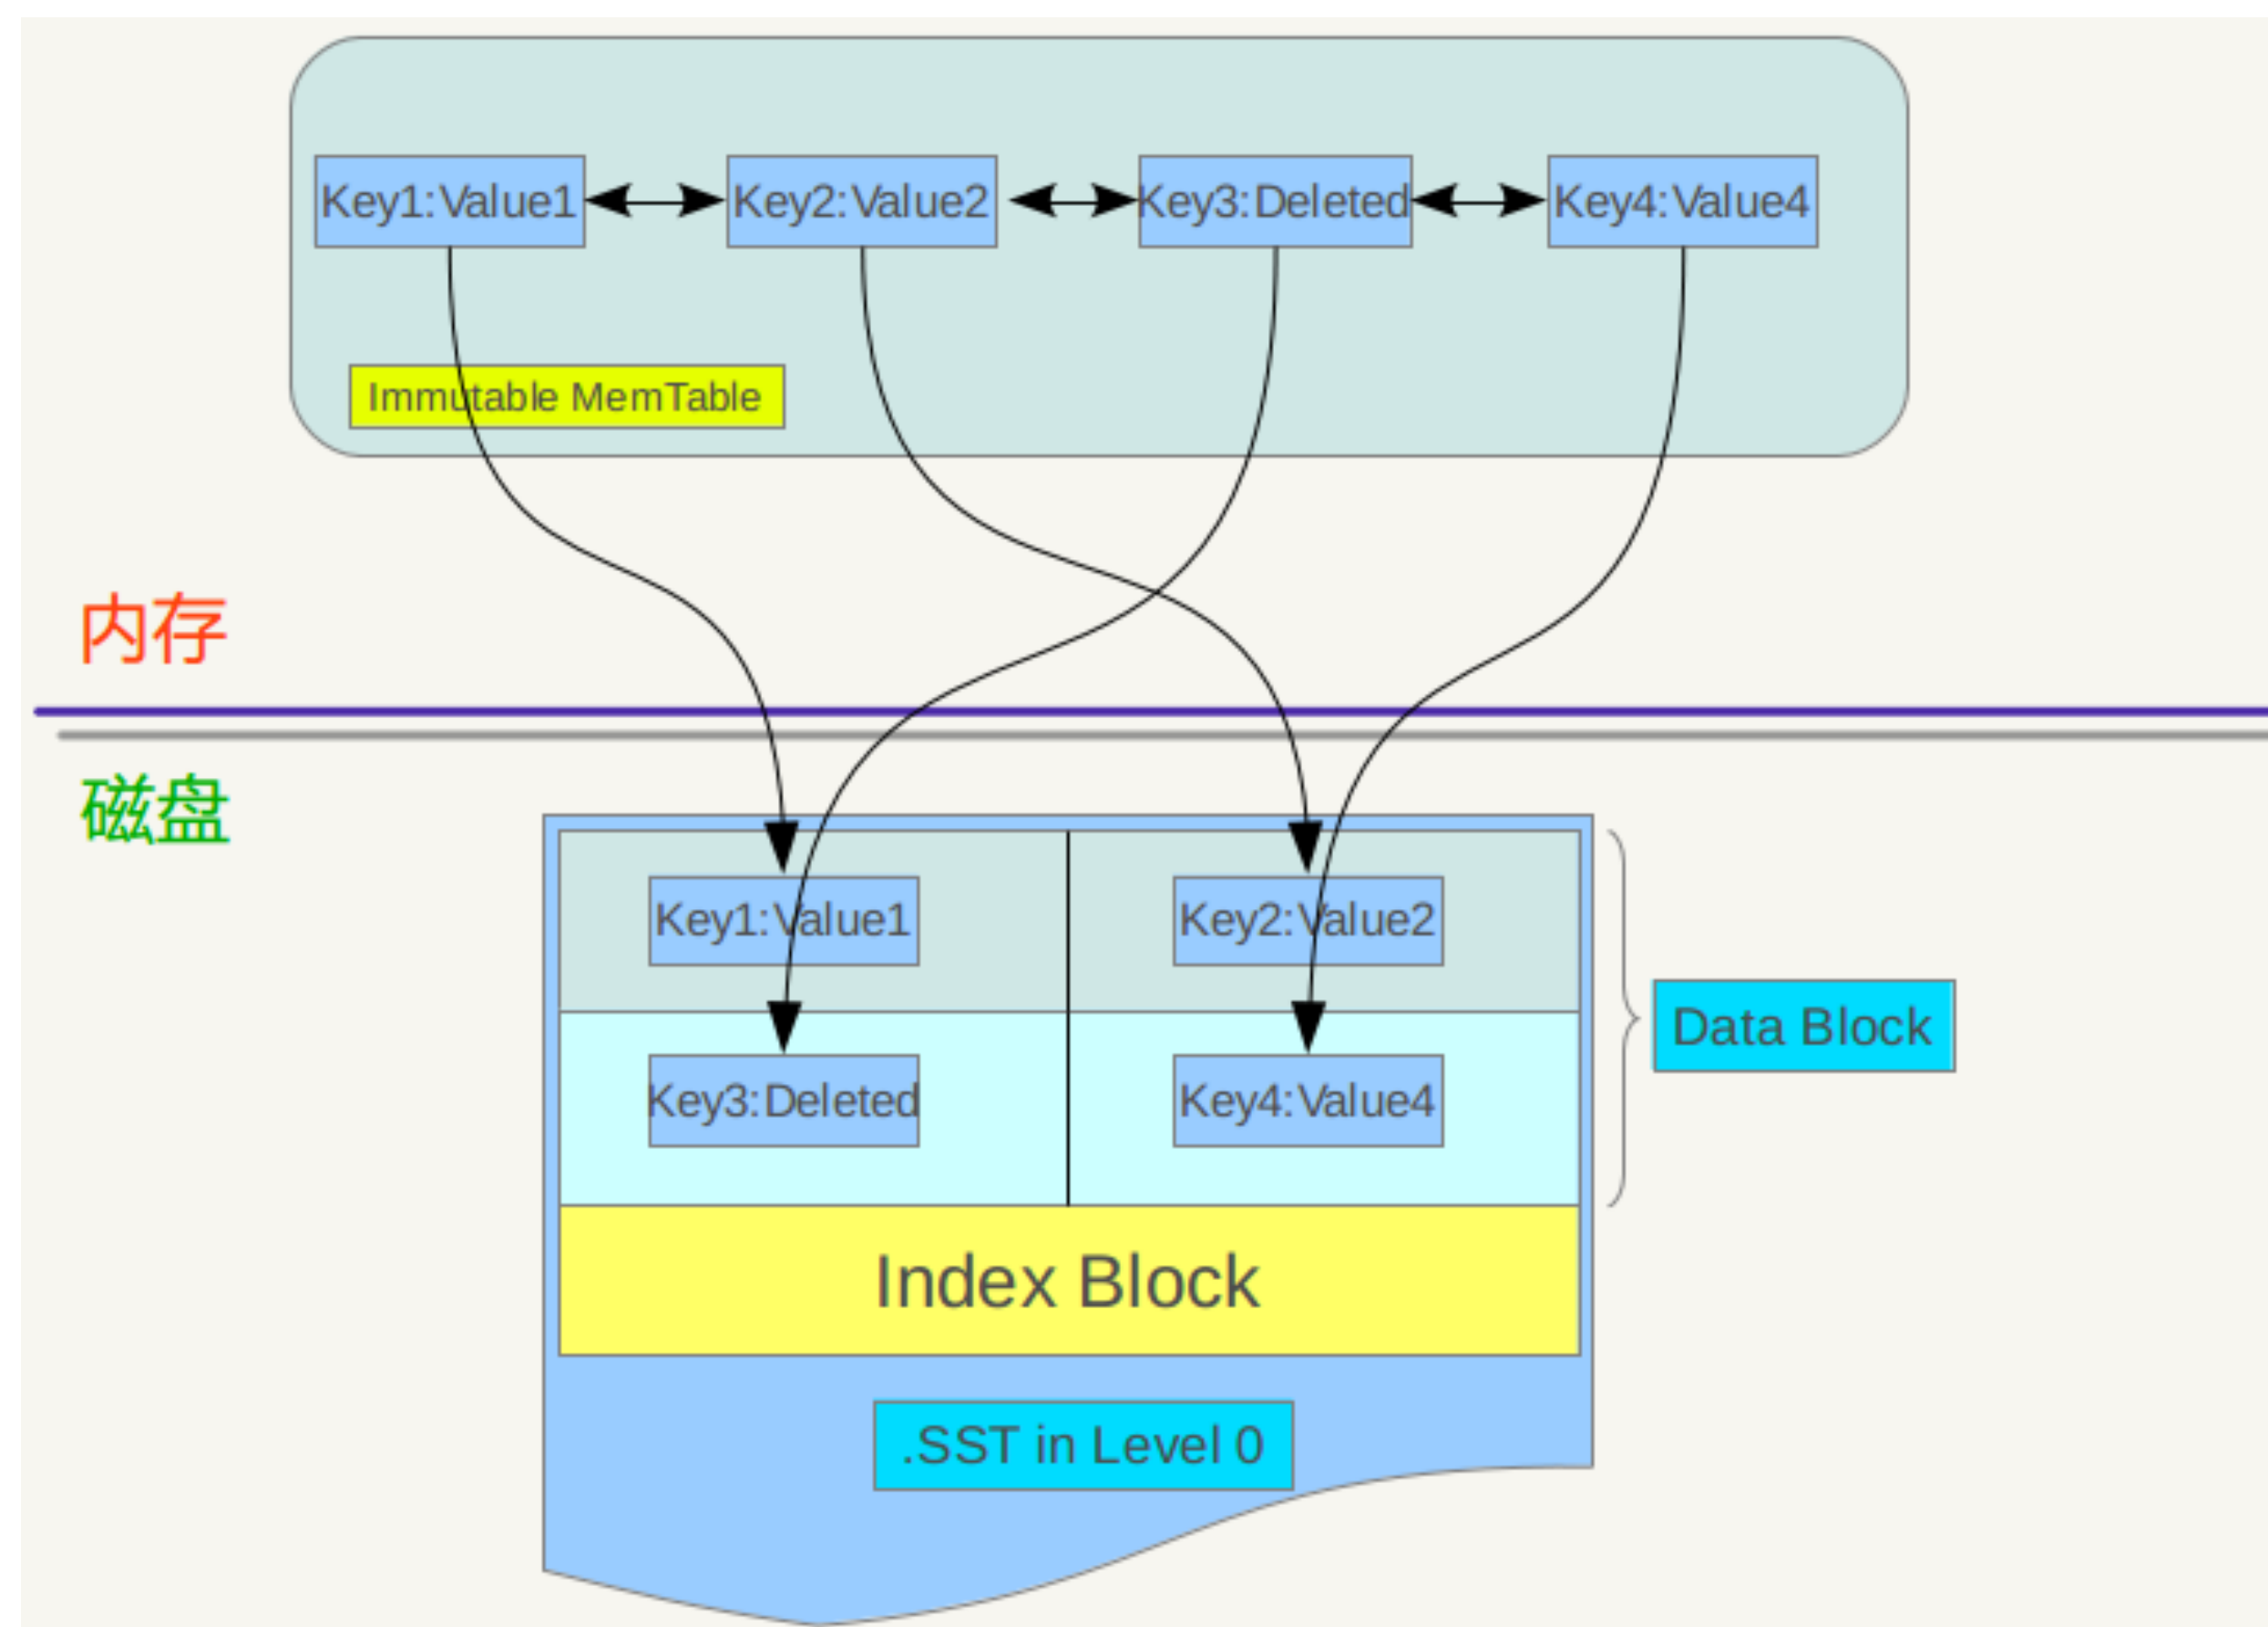


Write in LevelDB

- As mentioned before, if MemTable is full, we flushed it to disk as SSTable. But what if SSTable is full?
- We need Compaction
 - minor Compaction: merge memtable with level 0 SSTable
 - major compaction: merge SSTables in different levels
 - full compaction: merge all SSTables

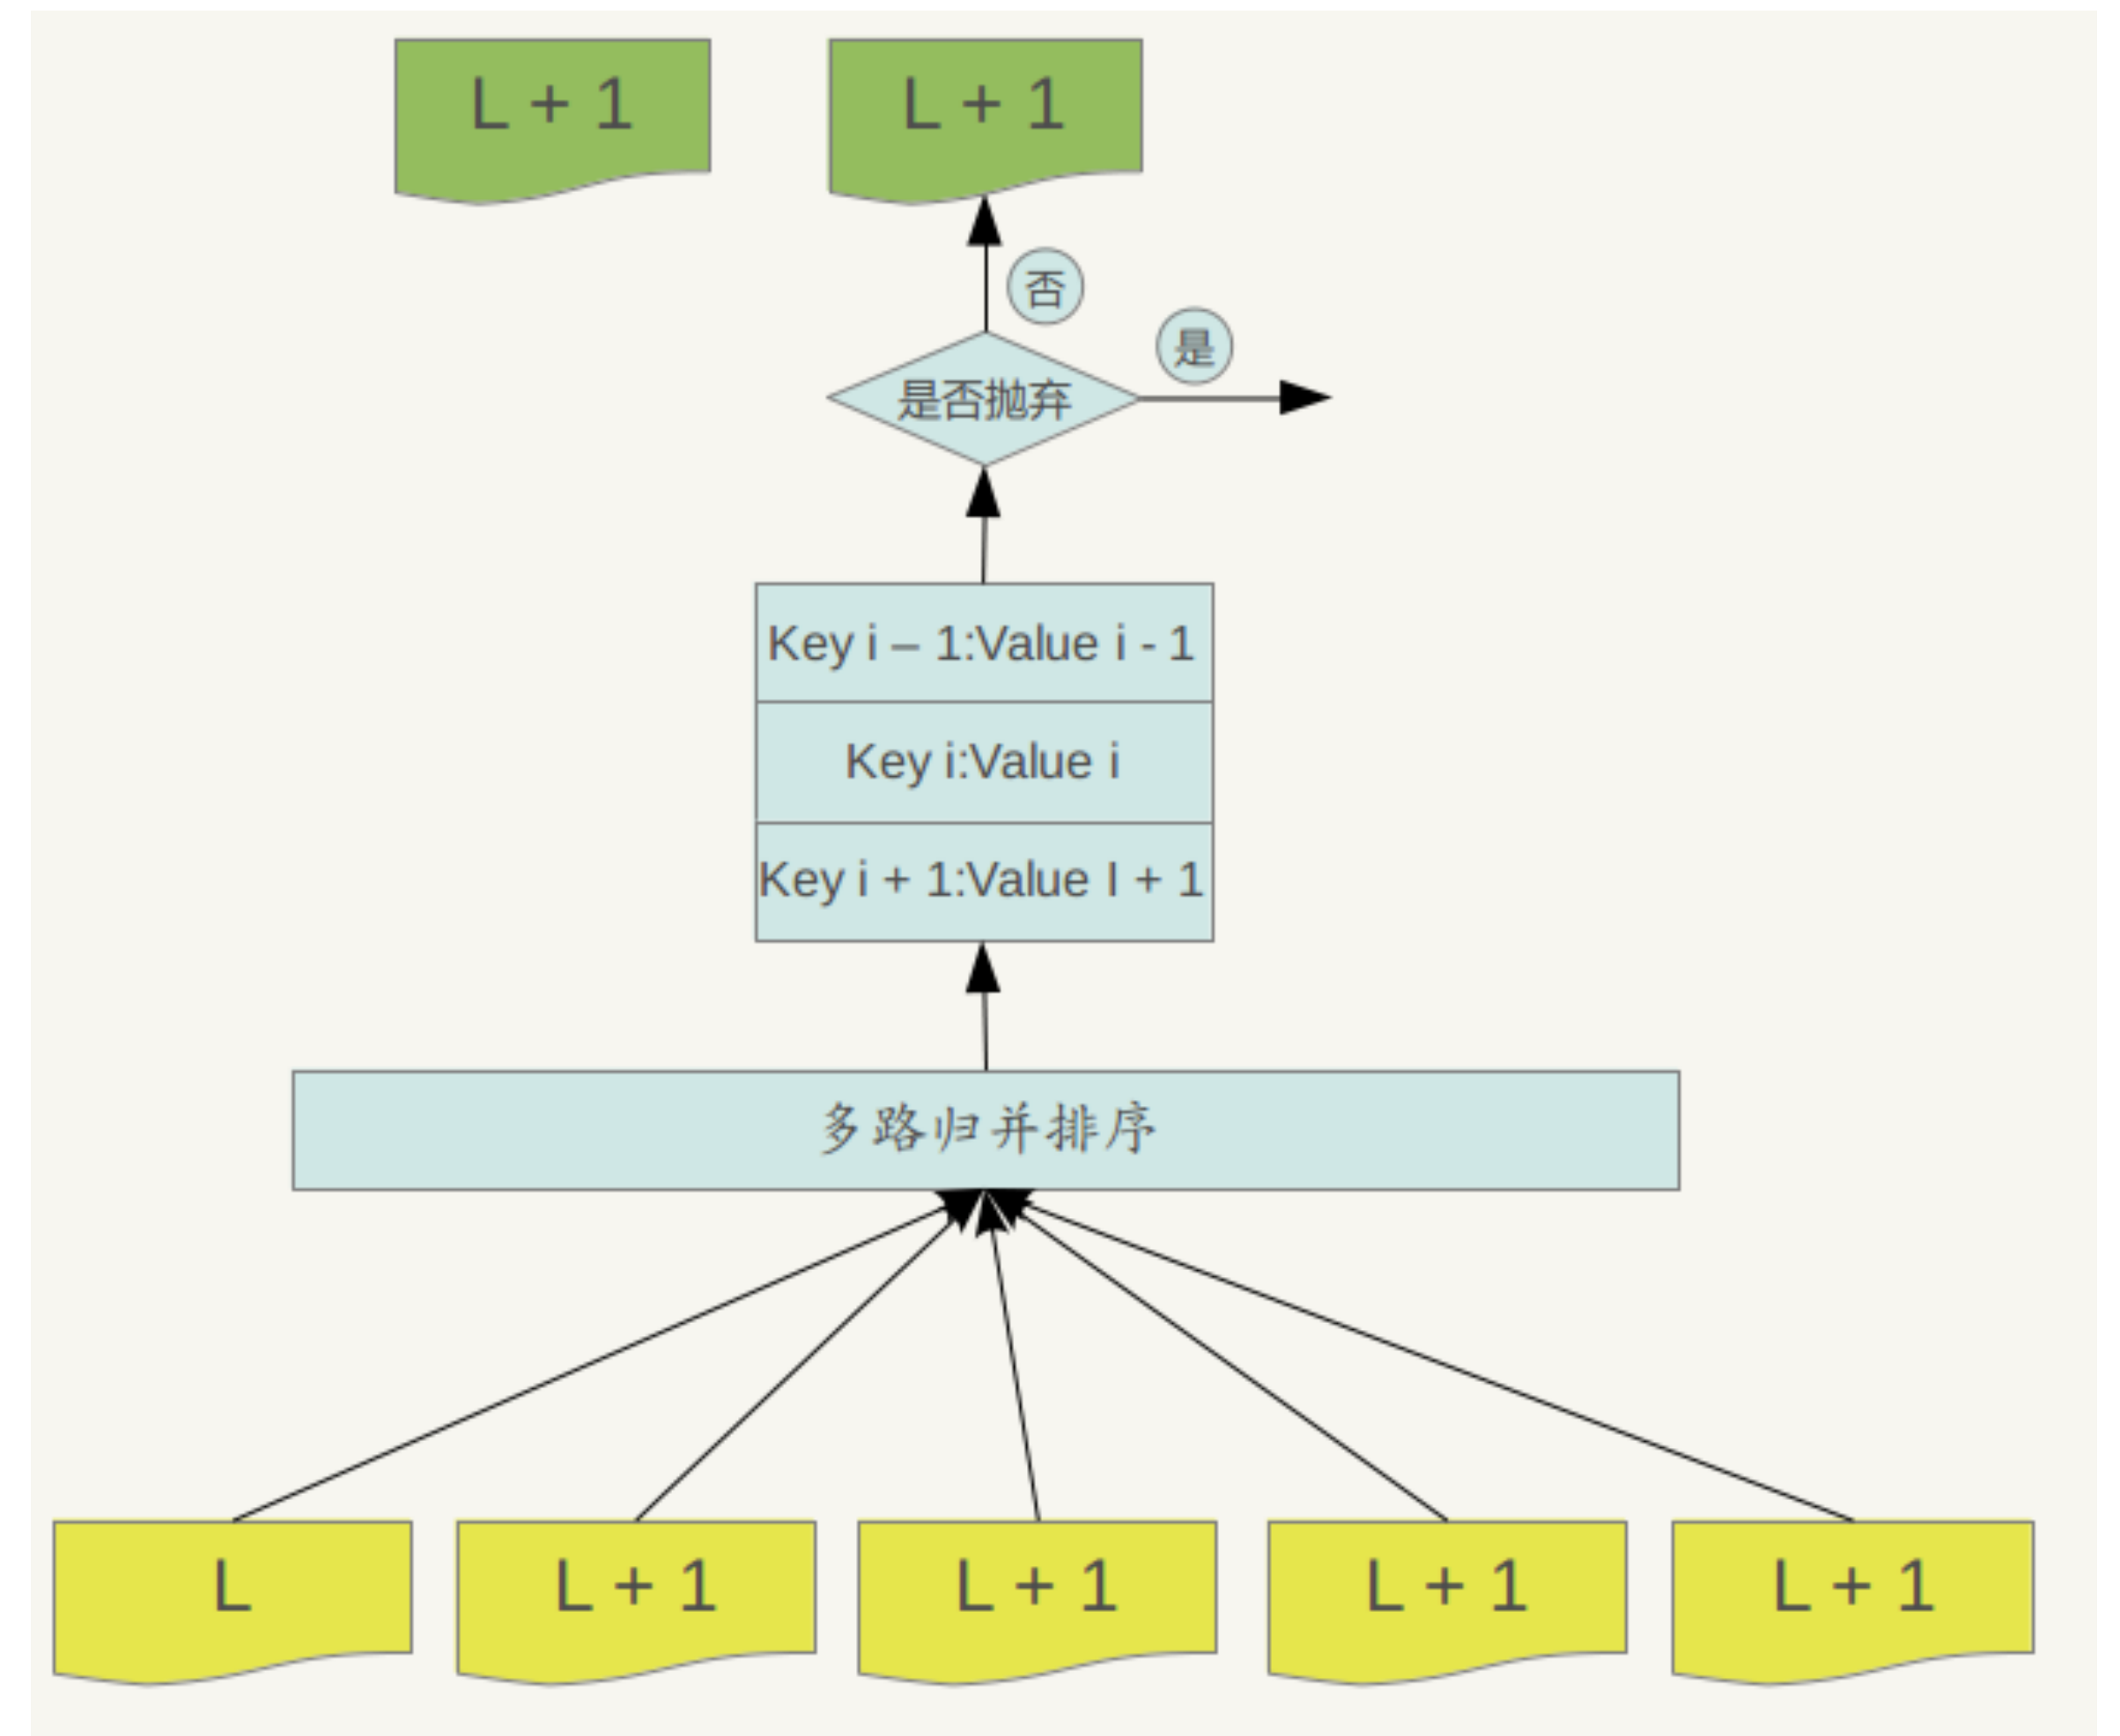


Minor Compaction




Major Compaction

- select a SSTable file in level L
- select SSTable files in level L+1 that have overlapped key ranges
- merge them together
- Note: For level 0, more than 1 SSTable files in level L are involved!




Still too slow: Cache

- Even if the data reside in level 0, we need to first check the Memtable and then the SSTable, which is slow.
- What if, the data are in level K?
- To avoid the case, we use a cache to maintain all hot data.




The Most Challenging Task: Multithreads

- It is OK for multi-threads to read together
- But if they write together, there may be problem.
- To guarantee the correctness, we need to lock the whole skiplist.
- That is too slow



What if you are the designer



One Solution (not the best)

- All writes are buffered in a queue and proceed one by one.
- The process who gets the access to the write operations also handles the read operations.

```
Status DBImpl::Write(const WriteOptions& options, WriteBatch*  
my_batch) {  
    // A begin  
    Writer w(&mutex_);  
    w.batch = my_batch;  
    w.sync = options.sync;  
    w.done = false;  
    // A end  
  
    // B begin  
    MutexLock l(&mutex_);  
    writers_.push_back(&w);  
    while (!w.done && &w != writers_.front()) {  
        w.cv.Wait();  
    }  
    if (w.done) {  
        return w.status;  
    }  
    // B end
```

Code Continues...

```
// May temporarily unlock and wait.
Status status = MakeRoomForWrite(my_batch == NULL);
uint64_t last_sequence = versions_->LastSequence();
Writer* last_writer = &w;
if (status.ok() && my_batch != NULL) { // NULL batch is for compactions
    WriteBatch* updates = BuildBatchGroup(&last_writer);
    WriteBatchInternal::SetSequence(updates, last_sequence + 1);
    last_sequence += WriteBatchInternal::Count(updates);

    // Add to log and apply to memtable. We can release the lock
    // during this phase since &w is currently responsible for logging
    // and protects against concurrent loggers and concurrent writes
    // into mem_.
    {
        mutex_.Unlock();
        status = log_->AddRecord(WriteBatchInternal::Contents(updates));
        bool sync_error = false;
        if (status.ok() && options.sync) {
            status = logfile_->Sync();
            if (!status.ok()) {
                sync_error = true;
            }
        }
    }
    if (status.ok()) {
        status = WriteBatchInternal::InsertInto(updates, mem_);
    }
}
```


And More


```
mutex_.Lock();
    if (sync_error) {
        RecordBackgroundError(status);
    }
}
if (updates == tmp_batch_) tmp_batch_->Clear();

versions_->SetLastSequence(last_sequence);
}

while (true) {
    Writer* ready = writers_.front();
    writers_.pop_front();
    if (ready != &w) {
        ready->status = status;
        ready->done = true;
        ready->cv.Signal();
    }
    if (ready == last_writer) break;
}

// Notify new head of write queue
if (!writers_.empty()) {
    writers_.front()->cv.Signal();
}

return status;
}
```



LevelDB Resources

- Website: <http://leveldb.org/>
- Source code: <https://github.com/google/leveldb/releases>
- papers:
 1. <https://static.googleusercontent.com/media/research.google.com/en//archive/bigtable-osdi06.pdf>
 2. <https://www.cs.cmu.edu/~ckingsf/bioinfo-lectures/skiplist.pdf>
 3. <http://citeseer.ist.psu.edu/viewdoc/download;jsessionid=6CA79DD1A90B3EFD3D62ACE5523B99E7?doi=10.1.1.127.9672&rep=rep1&type=pdf>

LEVELDB Opensourced

- RocksDB: Facebook version, <https://rocksdb.org>
- HBase: <https://hbase.apache.org>

Homework: deploy Hbase in the Cloud and test