

Lab 0: Linux内核重建及添加系统调用

姓名：韩耕诗

学号：3170103236

Part 1 Linux内核重建

1.1 Steps

查询当前版本 cat /proc/version, 当前内核为5.0.0版本.

```
cat /proc/version
```

```
hgs@hgs-vm:~$ cat /proc/version
Linux version 5.0.0-23-generic (bulld@lgw01-amd64-030) (gcc version 7.4.0 (Ubuntu 7.4.0-1ubuntu1~18.04.1)) #24~18.04.1-Ubuntu SMP Mon Jul 29 16:12:28 UTC 2019
hgs@hgs-vm:~$
```

查找源代码

```
sudo apt-cache search linux-source
```

```
hgs@hgs-vm:~$ sudo apt-cache search linux-source
[sudo] hgs 的密码:
linux-source - Linux kernel source with Ubuntu patches
linux-source-4.15.0 - Linux kernel source for version 4.15.0 with Ubuntu patches
linux-source-4.18.0 - Linux kernel source for version 4.18.0 with Ubuntu patches
linux-source-5.0.0 - Linux kernel source for version 5.0.0 with Ubuntu patches
```

下载5.0.0版本的内核

```
sudo apt-get install linux-source-5.0.0
```

```

hgs@hgs-vm:~$ sudo apt-get install linux-source-5.0.0
正在读取软件包列表... 完成
正在分析软件包的依赖关系树
正在读取状态信息... 完成
将会同时安装下列软件：
gcc gcc-7 libasan4 libatomic1 libc-dev-bin libc6-dev libcilkrts5
libgcc-7-dev libitm1 liblsan0 libmpx2 libquadmath0 libtsan0 libubsan0
linux-libc-dev make manpages-dev
建议安装：
gcc-multilib autoconf automake libtool flex bison gcc-doc gcc-7-multilib
gcc-7-doc gcc-7-locales libgcc1-dbg libgomp1-dbg libitm1-dbg libatomic1-dbg
libasan4-dbg liblsan0-dbg libtsan0-dbg libubsan0-dbg libcilkrts5-dbg
libmpx2-dbg libquadmath0-dbg glibc-doc libncurses-dev | ncurses-dev
kernel-package libqt3-dev make-doc
下列【新】软件包将被安装：
gcc gcc-7 libasan4 libatomic1 libc-dev-bin libc6-dev libcilkrts5
libgcc-7-dev libitm1 liblsan0 libmpx2 libquadmath0 libtsan0 libubsan0
linux-libc-dev linux-source-5.0.0 make manpages-dev
升级了 0 个软件包，新安装了 18 个软件包，要卸载 0 个软件包，有 140 个软件包未被
升级。
需要下载 149 MB 的归档。
解压缩后会消耗 223 MB 的额外空间。
您希望继续执行吗？ [Y/n] y
获取:1 http://cn.archive.ubuntu.com/ubuntu bionic-updates/main amd64 libitm1 am
d64 8.3.0-6ubuntu1~18.04.1 [28.0 kB]

```

部署内核源代码 将源代码解压到build_kernel/linux-source-5.0.0

```

cd
mkdir build_kernel
cd build_kernel
cp /usr/src/linux-source-4.13.0.tar.bz2 .
tar jxvf linux-source-4.13.0.tar.bz2

```

配置内核

```

cp /usr/src/linux-headers-5.0.0-23-generic/.config build_kernel/linux-source-
5.0.0
cd build_kernel/linux-source-4.13.0
make menuconfig

```

make menuconfig菜单出现后依次选择load, OK, save, OK, exit, exit

```

hgs@hgs-vm:~$ cp /usr/src/linux-headers-5.0.0-23-generic/.config build_kernel/linux-sou
rce-5.0.0
hgs@hgs-vm:~$ cd build_kernel/linux-source-5.0.0
hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$ make menuconfig
scripts/kconfig/mconf Kconfig

*** End of the configuration.
*** Execute 'make' to start the build or try 'make help'.

hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$

```

编译内核和模块，安装，生成启动文件

```

make -j4
sudo make modules_install -j 4
sudo make install -j 4
sudo update-initramfs -c -k 5.0.21
sudo update-grub

```

```

hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$ make -j4
scripts/kconfig/conf --synconfig Kconfig
DESCEND objtool
CALL scripts/checksyscalls.sh
CHK include/generated/compile.h
CC [M] net/bridge/netfilter/ebt_snat.o
CC [M] net/bridge/netfilter/ebt_log.o
CC [M] net/bridge/netfilter/ebt_nflog.o
DEPMOD 5.0.21
hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$ sudo make install -j 4
sh ./arch/x86/boot/install.sh 5.0.21 arch/x86/boot/bzImage \
    System.map "/boot"
run-parts: executing /etc/kernel/postinst.d/apt-auto-removal 5.0.21 /boot/vmlinuz-5.0.21
run-parts: executing /etc/kernel/postinst.d/initramfs-tools 5.0.21 /boot/vmlinuz-5.0.21
update-initramfs: Generating /boot/initrd.img-5.0.21
run-parts: executing /etc/kernel/postinst.d/unattended-upgrades 5.0.21 /boot/vmlinuz-5.0.21
run-parts: executing /etc/kernel/postinst.d/update-notifier 5.0.21 /boot/vmlinuz-5.0.21
run-parts: executing /etc/kernel/postinst.d/zz-update-grub 5.0.21 /boot/vmlinuz-5.0.21
Sourcing file `/etc/default/grub'
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-5.0.21
Found initrd image: /boot/initrd.img-5.0.21
Found linux image: /boot/vmlinuz-5.0.0-23-generic
Found initrd image: /boot/initrd.img-5.0.0-23-generic
Found memtest86+ image: /boot/memtest86+.elf
Found memtest86+ image: /boot/memtest86+.bin
done
hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$

```

```

done
hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$ ls -l /lib/modules/
总用量 8
drwxr-xr-x 5 root root 4096 8月 6 03:05 5.0.0-23-generic
drwxr-xr-x 3 root root 4096 9月 25 11:59 5.0.21
hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$ sudo update-initramfs -c -k 5.0.21
1
update-initramfs: Generating /boot/initrd.img-5.0.21
hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$

```

```

hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$ sudo update-grub
Sourcing file `/etc/default/grub'
Generating grub configuration file ...
Found linux image: /boot/vmlinuz-5.0.21
Found initrd image: /boot/initrd.img-5.0.21
Found linux image: /boot/vmlinuz-5.0.0-23-generic
Found initrd image: /boot/initrd.img-5.0.0-23-generic
Found memtest86+ image: /boot/memtest86+.elf
Found memtest86+ image: /boot/memtest86+.bin
done
hgs@hgs-vm:~/build_kernel/linux-source-5.0.0$

```

重启系统后验证

```

uname -a
dmesg | more

```

```

hgs@hgs-vm:~$ uname -a
Linux hgs-vm 5.0.21 #1 SMP Wed Sep 25 03:06:48 CST 2019 x86_64 x86_64 x86_64 GNU/Linux
hgs@hgs-vm:~$

```



```
hgs@hgs-vm: ~
文件(F) 编辑(E) 查看(V) 搜索(S) 终端(T) 帮助(H)
hgs@hgs-vm:~$ dmesg | more
[ 0.000000] Linux version 5.0.21 (hgs@hgs-vm) (gcc version 7.4.0 (Ubuntu 7.4
.0-1ubuntu1-18.04.1)) #1 SMP Wed Sep 25 03:06:48 CST 2019 (Ubuntu 5.0.0-23.24~1
8.04.1-generic 5.0.15)
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-5.0.21 root=UUID=55b11b64
-5404-45d3-bd41-e172402f93d0 ro quiet splash
[ 0.000000] KERNEL supported cpus:
[ 0.000000] Intel GenuineIntel
[ 0.000000] AMD AuthenticAMD
[ 0.000000] Hygon HygonGenuine
[ 0.000000] Centaur CentaurHauls
[ 0.000000] Disabled fast string operations
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x001: 'x87 floating point reg
isters'
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x002: 'SSE registers'
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x004: 'AVX registers'
[ 0.000000] x86/fpu: xstate_offset[2]: 576, xstate_sizes[2]: 256
[ 0.000000] x86/fpu: Enabled xstate features 0x7, context size is 832 bytes,
using 'standard' format.
[ 0.000000] BIOS-provided physical RAM map:
[ 0.000000] BIOS-e820: [mem 0x0000000000000000-0x00000000000009e7ff] usable
[ 0.000000] BIOS-e820: [mem 0x00000000000009e800-0x00000000000009ffff] reserved
[ 0.000000] BIOS-e820: [mem 0x0000000000000dc000-0x0000000000000fffff] reserved
[ 0.000000] BIOS-e820: [mem 0x000000000000100000-0x0000000000007fedffff] usable
[ 0.000000] BIOS-e820: [mem 0x0000000007fee00000-0x0000000007fefeffff] ACPI data
[ 0.000000] BIOS-e820: [mem 0x0000000007feff0000-0x0000000007feffffff] ACPI NVS
[ 0.000000] BIOS-e820: [mem 0x0000000007ff000000-0x0000000007ffffffffff] usable
[ 0.000000] BIOS-e820: [mem 0x000000000f00000000-0x000000000f7ffffffffff] reserved
```

1.2 Questions

1.提交你编译后的内核 dmesg | more 命令运行结果的截图，需要能明确显示出来你编译内核的时间和机器名，用户名（图片加框）

```
hgs@hgs-vm: ~
文件(F) 编辑(E) 查看(V) 搜索(S) 终端(T) 帮助(H)
hgs@hgs-vm:~$ dmesg | more
[ 0.000000] Linux version 5.0.21 (hgs@hgs-vm) (gcc version 7.4.0 (Ubuntu 7.4
.0-1ubuntu1-18.04.1)) #1 SMP Wed Sep 25 03:06:48 CST 2019 (Ubuntu 5.0.0-23.24~1
8.04.1-generic 5.0.15)
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-5.0.21 root=UUID=55b11b64
-5404-45d3-bd41-e172402f93d0 ro quiet splash
[ 0.000000] KERNEL supported cpus:
[ 0.000000] Intel GenuineIntel
[ 0.000000] AMD AuthenticAMD
[ 0.000000] Hygon HygonGenuine
[ 0.000000] Centaur CentaurHauls
[ 0.000000] Disabled fast string operations
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x001: 'x87 floating point reg
isters'
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x002: 'SSE registers'
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x004: 'AVX registers'
[ 0.000000] x86/fpu: xstate_offset[2]: 576, xstate_sizes[2]: 256
[ 0.000000] x86/fpu: Enabled xstate features 0x7, context size is 832 bytes,
using 'standard' format.
[ 0.000000] BIOS-provided physical RAM map:
[ 0.000000] BIOS-e820: [mem 0x0000000000000000-0x00000000000009e7ff] usable
[ 0.000000] BIOS-e820: [mem 0x00000000000009e800-0x00000000000009ffff] reserved
[ 0.000000] BIOS-e820: [mem 0x0000000000000dc000-0x0000000000000fffff] reserved
[ 0.000000] BIOS-e820: [mem 0x000000000000100000-0x0000000000007fedffff] usable
[ 0.000000] BIOS-e820: [mem 0x0000000007fee00000-0x0000000007fefeffff] ACPI data
[ 0.000000] BIOS-e820: [mem 0x0000000007feff0000-0x0000000007feffffff] ACPI NVS
[ 0.000000] BIOS-e820: [mem 0x0000000007ff000000-0x0000000007ffffffffff] usable
[ 0.000000] BIOS-e820: [mem 0x000000000f00000000-0x000000000f7ffffffffff] reserved
```

2.Linux 内核目录下有一个.config 文件，请说明这个文件的作用？

.config文件是内核配置文件，顶层Makefile会使用这个.config文件来构建内核。大多数内核软件模块也通过.config文件间接地读取配置内容。

3.在Linux内核代码树中，很多子目录Makefile 文件和 Kconfig文件，请分别解释这两个文件的作用？

每个Kconfig分别描述了所属目录源文档相关的内核配置菜单，在执行内核配置make menuconfig时，会从Kconfig中读出菜单，用户选择后保存到.config的内核配置文档中。

在内核编译时，主Makefile调用这个基于Kconfig生成的.config，就知道了用户的选择。

要想添加新的驱动到内核的源码中，就要修改Kconfig，这样就能够选择这个驱动。

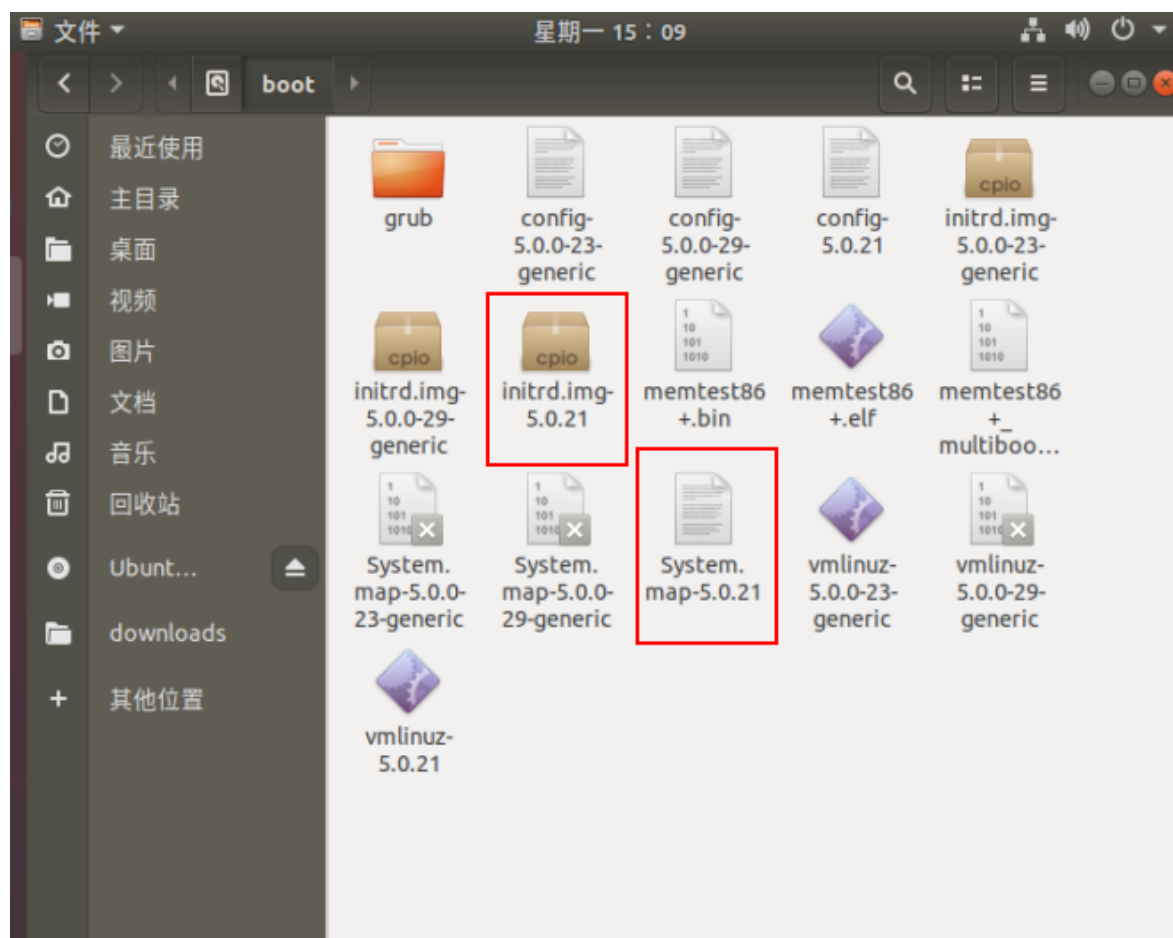
如果想使新的驱动被编译，则要修改Makefile文件。

4.浏览/boot 目录，你一定发现了 System.map-4.13.16 文件，以及 initrd.img-4.13.0 文件。这两个文件分别起什么作用？你能否设计一个实验来验证你的判断？

initrd 的含义是initial RAM disk，就是启动系统所需加载的虚拟磁盘。在系统启动过程中，kernel、initrd和system module是依次加载的。initrd包含一部分内核模块，主要是一些关键的外部硬件，如 SATA、SCSI和USB等外设。如果initrd不存在或者失败会影响系统启动。

System.map文件是对应内核的符号映射表，顾名思义就是将内核中的符号（包括变量和函数）和它的地址联系起来的一个列表，是所有符号及其对应地址的一个列表。

设计一个实验：将这两个文件复制到其他文件夹后，在boot文件夹下删除这两个文件，重启系统。



删除这两个文件后重启系统出现错误。

```
BusyBox v1.27.2 (Ubuntu 1:1.27.2-2ubuntu3.2) built-in shell (ash)
Enter 'help' for a list of built-in commands.

(initramfs)
```

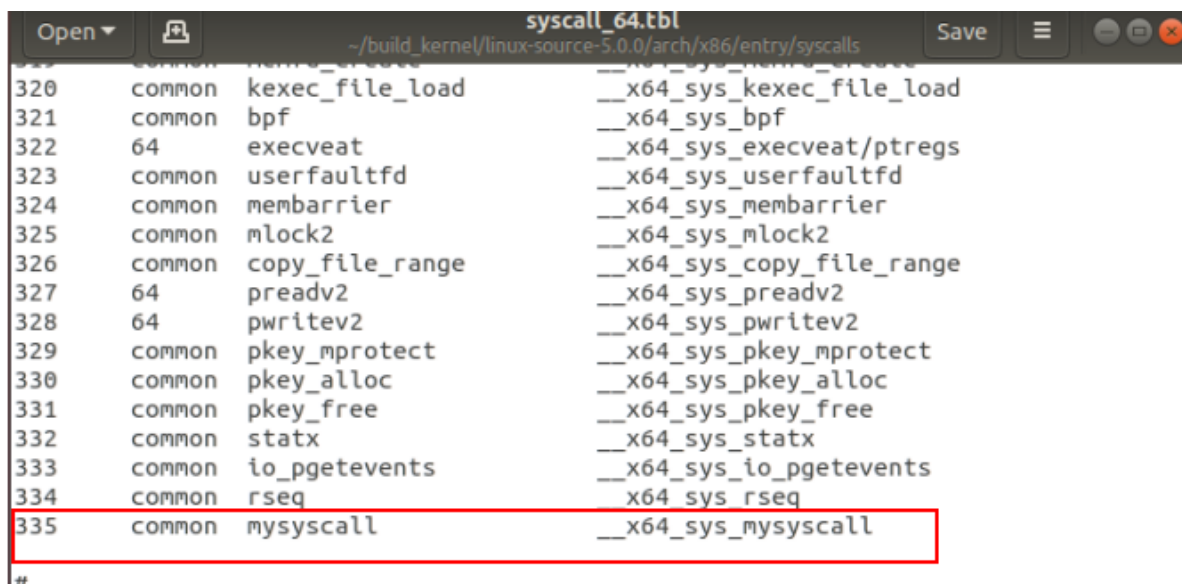
开机高级选项进入之前的内核版本5.0.23-generic，将这两个文件恢复，重启成功。

Part 2 添加系统调用

2.1 Steps

2.1.1 在系统调用表中添加相应表项

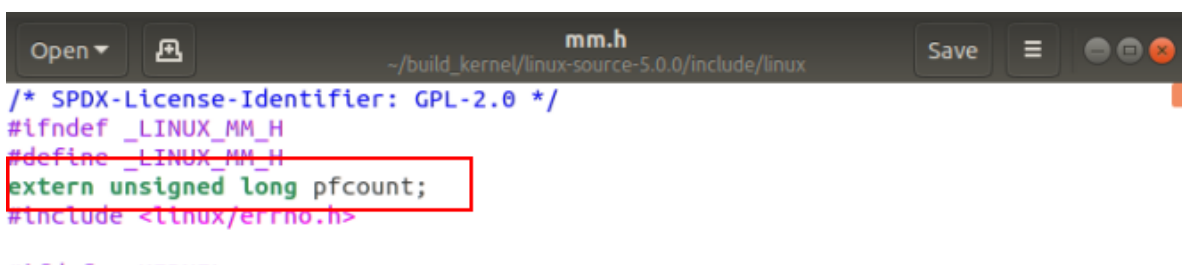
找到arch/x86/entry/syscalls/syscall_64.tbl，在common中添加一条



| Common | Number | Common Name | 64-bit Implementation |
|--------|--------|-----------------|---------------------------|
| common | 320 | kexec_file_load | __x64_sys_kexec_file_load |
| common | 321 | bpf | __x64_sys_bpf |
| 64 | 322 | execveat | __x64_sys_execveat/ptregs |
| common | 323 | userfaultfd | __x64_sys_userfaultfd |
| common | 324 | membarrier | __x64_sys_membarrier |
| common | 325 | mlock2 | __x64_sys_mlock2 |
| common | 326 | copy_file_range | __x64_sys_copy_file_range |
| 64 | 327 | preadv2 | __x64_sys_preadv2 |
| 64 | 328 | pwritev2 | __x64_sys_pwritev2 |
| common | 329 | pkey_mprotect | __x64_sys_pkey_mprotect |
| common | 330 | pkey_alloc | __x64_sys_pkey_alloc |
| common | 331 | pkey_free | __x64_sys_pkey_free |
| common | 332 | statx | __x64_sys_statx |
| common | 333 | io_pgetevents | __x64_sys_io_pgetevents |
| common | 334 | rseq | __x64_sys_rseq |
| common | 335 | mysyscall | __x64_sys_mysyscall |

2.1.2 修改统计系统缺页次数和进程缺页次数的内核代码

先在 include/linux/mm.h 文件中声明变量 pfcount



```
/* SPDX-License-Identifier: GPL-2.0 */
#ifndef _LINUX_MM_H
#define _LINUX_MM_H
extern unsigned long pfcount;
#include <linux/errno.h>
```

在进程 task_struct 中增加成员 pf，在 include/linux/sched.h 文件中的 task_struct 结构中添加 pf 字段

```
Open ▾  sched.h  Save  ⋮  ⌵  ✖
~/build_kernel/linux-source-5.0.0/include/linux

#ifdef CONFIG_NO_HZ_FULL
    atomic_t          tick_dep_mask;
#endif

    /* Context switch counts: */
    unsigned long      nvcs;
    unsigned long      nivcs;

    unsigned long pf;

    /* Monotonic time in nsecs: */
    u64                start_time;
```

统计当前进程缺页次数需要在创建进程是需要将进程控制块中的pf设置为0，在进程创建过程中，子进程会把父进程的进程控制块复制一份，实现该复制过程的函数是kernel/fork.c文件中的dup_task_struct()函数，修改该函数将子进程的pf设置成0。

```
Open ▾  fork.c  Save  ⋮  ⌵  ✖
~/build_kernel/linux-source-5.0.0/kernel

    *stackend = STACK_END_MAGIC; /* for overflow detection */
}

static struct task_struct *dup_task_struct(struct task_struct *orig, int node)
{
    struct task_struct *tsk;
    unsigned long *stack;
    struct vm_struct *stack_vm_area __maybe_unused;
    int err;

    if (node == NUMA_NO_NODE)
        node = tsk_fork_get_node(orig);
    tsk = alloc_task_struct_node(node);
    if (!tsk)
        return NULL;

    stack = alloc_thread_stack_node(tsk, node);
    if (!stack)
        goto free_tsk;

    if (memcg_charge_kernel_stack(tsk))
        goto free_stack;

    tsk->pf = 0;

    stack_vm_area = task_stack_vm_area(tsk);
```

在arch/x86/mm/fault.c文件中定义变量pfcount，并修改arch/x86/mm/fault.c中do_page_fault()函数。每次产生缺页中断，do_page_fault()函数会被调用，pfcount变量值递增1，记录系统产生缺页次数，current->pf值递增1，记录当前进程产生缺页次数。


```
hgs2@hgs2-vm:~$ dmesg | more
[ 0.000000] Linux version 5.0.21 (hgs2@hgs2-vm) (gcc version 7.4.0 (Ubuntu 7
.4.0-1ubuntu1~18.04.1)) #2 SMP Tue Oct 1 16:38:53 CST 2019 (Ubuntu 5.0.0-23.24~
18.04.1-generic 5.0.15)
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-5.0.21 root=UUID=9464a22b
-0cf2-446c-99f7-06f0c7219c6e ro quiet splash
[ 0.000000] KERNEL supported cpus:
[ 0.000000] Intel GenuineIntel
[ 0.000000] AMD AuthenticAMD
[ 0.000000] AMD AuthenticAMD
```

2.1.5 编写用户态程序

编写test.c

```
#include <linux/unistd.h>
#include <sys/syscall.h>
#define __NR_mysyscall 335
int main()
{
    syscall(__NR_mysyscall);
}
```

编译运行test, 并且用dmesg验证

```
gcc -o test test.c
./test
dmesg
```

```
hgs2@hgs2-vm:~$ ./test
hgs2@hgs2-vm:~$ dmesg
[ 0.000000] Linux version 5.0.21 (hgs2@hgs2-vm) (gcc version 7.4.0 (Ubuntu 7
.4.0-1ubuntu1~18.04.1)) #2 SMP Tue Oct 1 16:38:53 CST 2019 (Ubuntu 5.0.0-23.24~
18.04.1-generic 5.0.15)
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-5.0.21 root=UUID=9464a22b
-0cf2-446c-99f7-06f0c7219c6e ro quiet splash
[ 0.000000] KERNEL supported cpus:
[ 0.000000] Intel GenuineIntel
[ 0.000000] AMD AuthenticAMD
[ 0.000000] Hygon HygonGenuine
[ 0.000000] Centaur CentaurHauls
[ 0.000000] Disabled fast string operations
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x001: 'x87 floating point reg
isters'
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x002: 'SSE registers'
[ 0.000000] x86/fpu: Supporting XSAVE feature 0x004: 'AVX registers'
[ 0.000000] x86/fpu: xstate offset[2]: 576 xstate sizes[2]: 256
[ 0.000000] x86/fpu: xstate offset[3]: 1280 xstate sizes[3]: 1280
[ 649.052205] current process - page fault count 15166
[ 1021.875965] current process - page fault count 15286
[ 1048.677554] current process - page fault count 15399
```

2.2 Questions

1.在 test.c 中添加打印整个系统 page fault 的变量值 (也就是 pfcount 的值)。上传你的 test.c 代码。

修改test.c以及sys.c

```
test.c
~/
Open Save
#include <linux/unistd.h>
#include <sys/syscall.h>
#include <stdio.h>
#define __NR_mysyscall 335

int main()
{
    unsigned long pfcoun;
    pfcoun = syscall(__NR_mysyscall);
    printf("system page fault count(pfcoun) = %ld \n",pfcoun);
}
```

```
sys.c
~/build_kernel/linux-source-5.0.0/kernel
Open Save
test.c sys.c
}
SYSCALL_DEFINE0(mysyscall)
{
    printk("current process - page fault count %ld \n", current->pf);
    return pfcoun;
}
/*
 * Accessing ->real_parent is not SMP-safe, it could
 * change from under us. However, we can use a stale
```

test.c源代码修改如下:

```
#include <linux/unistd.h>
#include <sys/syscall.h>
#include <stdio.h>
#define __NR_mysyscall 335
int main()
{
    unsigned long pfcoun;
    pfcoun = syscall(__NR_mysyscall);
    printf("system page fault count(pfcoun) = %ld \n",pfcoun);
}
```

重新编译内核

```
make -j4
sudo make modules_install -j 4
sudo make install -j 4
sudo update-initramfs -c -k 5.0.21
sudo update-grub
```

重启系统, 用 dmesg | more 验证内核编译成功

```
hgs2@hgs2-vm: ~
File Edit View Search Terminal Help
hgs2@hgs2-vm:~$ dmesg | more
[ 0.000000] Linux version 5.0.21 (hgs2@hgs2-vm) (gcc version 7.4.0 (Ubuntu 7
.4.0-1ubuntu1~18.04.1)) #4 SMP Tue Oct 1 19:01:23 CST 2019 (Ubuntu 5.0.0-23.24~
18.04.1-generic 5.0.15)
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-5.0.21 root=UUID=9464a22b
-0cf2-446c-99f7-06f0c7219c6e ro quiet splash
[ 0.000000] KERNEL supported cpus:
[ 0.000000] Intel GenuineIntel
[ 0.000000] AMD AuthenticAMD
[ 0.000000] Hygon HygonCseries
```

重新用gcc编译test程序

```
hgs2@hgs2-vm: ~
File Edit View Search Terminal Help
hgs2@hgs2-vm:~$ gcc -o test test.c
test.c: In function 'main':
test.c:9:12: warning: implicit declaration of function 'syscall'; did you mean
'sscanf'? [-Wimplicit-function-declaration]
    pfcount = syscall(__NR_mysyscall);
               ^~~~~~
               sscanf
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 646067
hgs2@hgs2-vm:~$
```

2.运行 test 程序后，dmesg 的截图证明你的系统调用添加成功，并且能在用户态被调用。

```
./test
dmesg
```

如下图成功运行test程序，dmesg末尾有调用日志。

```
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 646067
hgs2@hgs2-vm:~$ dmesg
[ 171.425853] current process - page fault count 14737
hgs2@hgs2-vm:~$
```

3.多次运行test程序，每次运行test后记录下系统缺页次数和当前进程缺页次数。

(后面第7题改正实验指导中存在的错误后，重新做了一遍第3题)

多次运行test以及dmesg

```
./test
dmesg
```

第一次运行test，系统缺页次数646067，当前进程缺页次数14737

```
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 646067
hgs2@hgs2-vm:~$ dmesg
[ 171.425853] current process - page fault count 14737
hgs2@hgs2-vm:~$
```

第二次运行test，系统缺页次数651051，当前进程缺页次数14838

```
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 651051
hgs2@hgs2-vm:~$ dmesg
[ 171.425853] current process - page fault count 14737
[ 353.914321] current process - page fault count 14838
hgs2@hgs2-vm:~$
```

第三次运行test，系统缺页次数651809，当前进程缺页次数14928

```
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 651809
hgs2@hgs2-vm:~$ dmesg
[ 171.425853] current process - page fault count 14737
[ 353.914321] current process - page fault count 14838
[ 438.484637] current process - page fault count 14928
hgs2@hgs2-vm:~$
```

第四次运行test，系统缺页次数652233，当前进程缺页次数15026

```
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 652233
hgs2@hgs2-vm:~$ dmesg
[ 171.425853] current process - page fault count 14737
[ 353.914321] current process - page fault count 14838
[ 438.484637] current process - page fault count 14928
[ 511.112083] current process - page fault count 15026
hgs2@hgs2-vm:~$
```

第五次运行test，系统缺页次数652703，当前进程缺页次数15118

```
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 652703
hgs2@hgs2-vm:~$ dmesg
[ 171.425853] current process - page fault count 14737
[ 353.914321] current process - page fault count 14838
[ 438.484637] current process - page fault count 14928
[ 511.112083] current process - page fault count 15026
[ 567.691433] current process - page fault count 15118
hgs2@hgs2-vm:~$
```

4.除了通过修改内核来添加一个系统调用外，还有其他的添加或修改一个系统调用的方法吗？如果有，请论述。

除了内核编译法来添加系统调用之外，也可以通过module进行内核添加来添加系统调用。这种方法是采用系统调用拦截的一种方式，改变某一个系统调用号对应的服务程序，变为自己编写的程序，从而相当于添加了系统调用。

5.对于一个操作系统而言，你认为修改系统调用的方法安全吗？请发表你的观点。

我认为对于一个操作系统而言，修改系统调用并不够安全。因为当修改系统调用时，有可能会对原来系统调用表中的其他调用进行修改，从而使得其名称发生变化或者是缺少对应编号的系统调用，可能会给调用正常系统调用接口的准确性和安全性造成威胁。

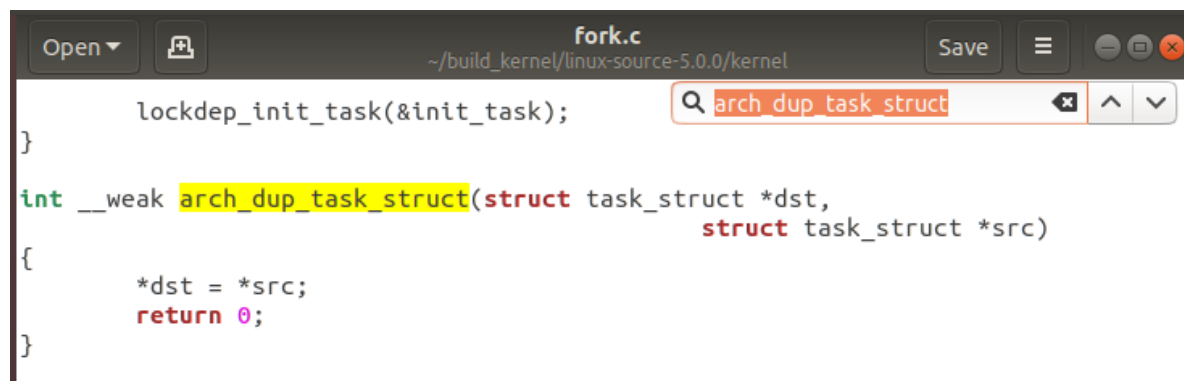
6.在实验过程中遇到了什么问题，你是如何解决的。

在实验过程中，我在添加完系统调用，重新编译内核后，重启系统失败，系统无法开机。因为在课程群里看到有同学出现了相同的问题是重装系统解决的，所以我也重新安装了Ubuntu，再次编译内核后再添加系统调用以及重新编译，这次就成功了。

7.在实验指导2_添加系统调用里有一个错误，希望大家能发现并写在作业中，作为第7个问题。 (hint: 子进程的pf如何保证从0开始计数?)

“tsk->pf = 0;”的位置不对。在kernel/fork.c文件中，“tsk->pf = 0”应该放在“err = arch_dup_task_struct (tsk, orig);”函数调用语句之后，因为这个函数里面有对tsk->pf造成改变。

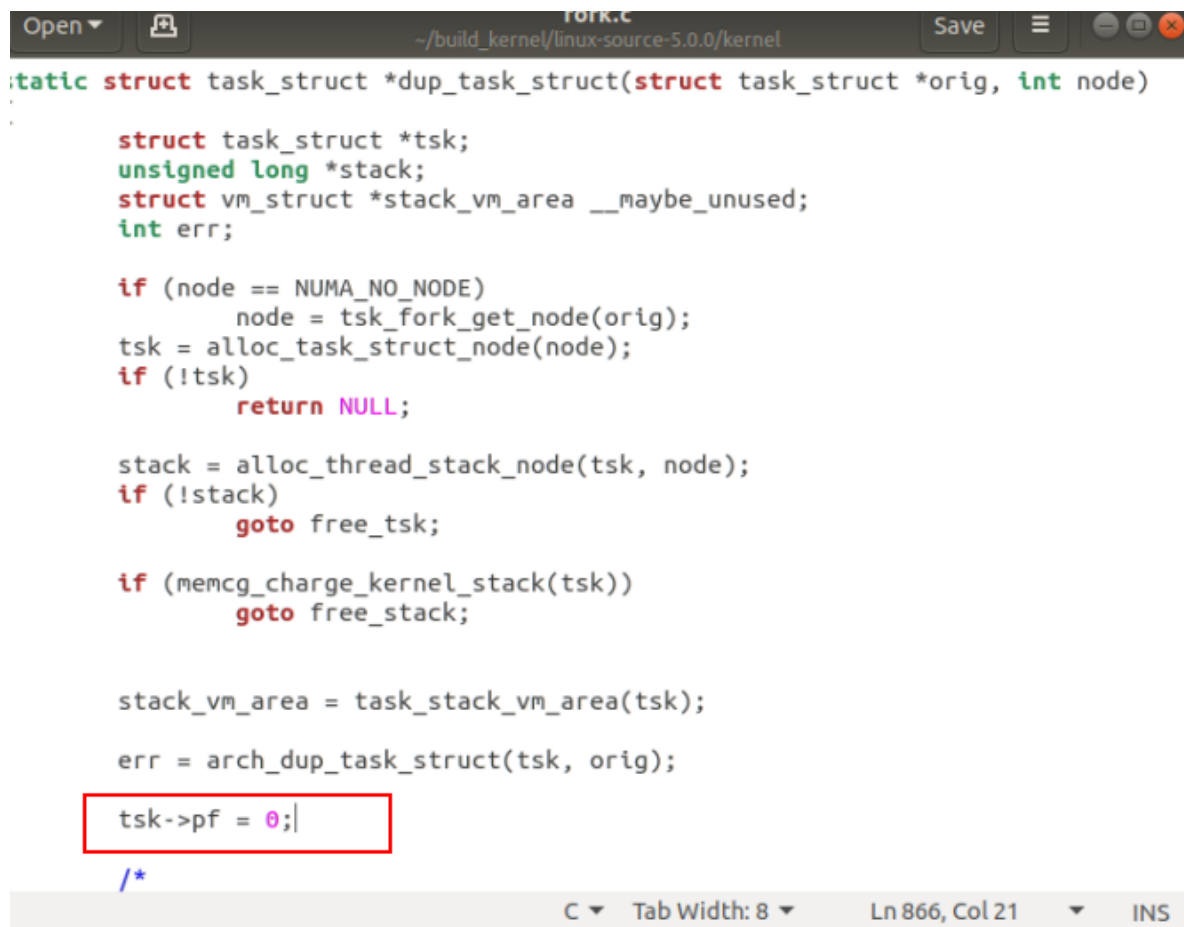
如下图是arch_dup_task_struct () 函数



The screenshot shows a code editor window titled 'fork.c' with the path '~/build_kernel/linux-source-5.0.0/kernel'. A search bar at the top right contains 'arch_dup_task_struct'. The function definition is as follows:

```
lockdep_init_task(&init_task);  
}  
  
int __weak arch_dup_task_struct(struct task_struct *dst,  
                               struct task_struct *src)  
{  
    *dst = *src;  
    return 0;  
}
```

经过改正后的kernel/fork.c中的dup_task_struct()函数如下



The screenshot shows the same code editor window with the 'dup_task_struct' function. The function definition is as follows:

```
static struct task_struct *dup_task_struct(struct task_struct *orig, int node)  
{  
    struct task_struct *tsk;  
    unsigned long *stack;  
    struct vm_struct *stack_vm_area __maybe_unused;  
    int err;  
  
    if (node == NUMA_NO_NODE)  
        node = tsk_fork_get_node(orig);  
    tsk = alloc_task_struct_node(node);  
    if (!tsk)  
        return NULL;  
  
    stack = alloc_thread_stack_node(tsk, node);  
    if (!stack)  
        goto free_tsk;  
  
    if (memcg_charge_kernel_stack(tsk))  
        goto free_stack;  
  
    stack_vm_area = task_stack_vm_area(tsk);  
    err = arch_dup_task_struct(tsk, orig);  
    tsk->pf = 0;  
    /*  
    */  
}
```

The line 'tsk->pf = 0;' is highlighted with a red box.

改正后重新编译内核

```
make -j4
sudo make modules_install -j 4
sudo make install -j 4
sudo update-initramfs -c -k 5.0.21
sudo update-grub
```

重启系统，用 dmesg | more 验证内核编译成功

```
hgs2@hgs2-vm: ~
File Edit View Search Terminal Help
hgs2@hgs2-vm:~$ dmesg | more
[ 0.000000] Linux version 5.0.21 (hgs2@hgs2-vm) (gcc version 7.4.0 (Ubuntu 7
.4.0-1ubuntu1~18.04.1)) #5 SMP Tue Oct 1 21:03:49 CST 2019 (Ubuntu 5.0.0-23.24~
18.04.1-generic 5.0.15)
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-5.0.21 root=UUID=9464a22b
-0cf2-446c-99f7-06f0c7219c6e ro quiet splash
[ 0.000000] KERNEL supported cpus:
[ 0.000000] Intel GenuineIntel
[ 0.000000] AMD AuthenticAMD
```

重新编译运行test程序以及dmesg

```
hgs2@hgs2-vm:~$ gcc -o test test.c
test.c: In function 'main':
test.c:9:12: warning: implicit declaration of function 'syscall'; did you mean
'sscanf'? [-Wimplicit-function-declaration]
    pfcount = syscall(__NR_mysyscall);
               ^~~~~~
               sscanf
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 629229
hgs2@hgs2-vm:~$ dmesg
[ 0.000000] Linux version 5.0.21 (hgs2@hgs2-vm) (gcc version 7.4.0 (Ubuntu 7
.4.0-1ubuntu1~18.04.1)) #5 SMP Tue Oct 1 21:03:49 CST 2019 (Ubuntu 5.0.0-23.24~
18.04.1-generic 5.0.15)
[ 0.000000] Command line: BOOT_IMAGE=/boot/vmlinuz-5.0.21 root=UUID=9464a22b
-0cf2-446c-99f7-06f0c7219c6e ro quiet splash
[ 0.000000] Extensible: kkr1_1991h
[ 237.711231] current process - page fault count 66
```

从当前进程的page fault count 来看，改正后数值是比较合理的

那么重新做第2题：多次运行test程序，每次运行test后记录下系统缺页次数和当前进程缺页次数

第1次运行test

```
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 629229
hgs2@hgs2-vm:~$ dmesg
[ 0.000000] Linux version 5.0.21 (hgs2@hgs2-vm) (gcc version 7.4.0 (Ubuntu 7
.4.0-1ubuntu1~18.04.1)) #5 SMP Tue Oct 1 21:03:49 CST 2019 (Ubuntu 5.0.0-23.24~
18.04.1-generic 5.0.15)
[ 0.000000] Extensible: kkr1_1991h
[ 237.711231] current process - page fault count 66
```

第2次运行test

```
[ 237.711231] current process - page fault count 66
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 638594
hgs2@hgs2-vm:~$ dmesg
[ 237.711231] current process - page fault count 66
[ 705.374079] current process - page fault count 68
```

第3次运行test

```
[ 705.374079] current process - page fault count 68
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 639550
hgs2@hgs2-vm:~$ dmesg
```

```
[ 237.711231] current process - page fault count 66
[ 705.374079] current process - page fault count 68
[ 786.670646] current process - page fault count 66
hgs2@hgs2-vm:~$
```

第4次运行test

```
[ 705.374079] current process - page fault count 68
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 642137
hgs2@hgs2-vm:~$ dmesg
```

```
[ 237.711231] current process - page fault count 66
[ 705.374079] current process - page fault count 68
[ 786.670646] current process - page fault count 66
[ 937.591141] current process - page fault count 67
hgs2@hgs2-vm:~$
```

第5次运行test

```
hgs2@hgs2-vm:~$ ./test
system page fault count(pfcount) = 642577
hgs2@hgs2-vm:~$ dmesg
```

```
[ 237.711231] current process - page fault count 66
[ 705.374079] current process - page fault count 68
[ 786.670646] current process - page fault count 66
[ 937.591141] current process - page fault count 67
[ 981.694806] current process - page fault count 66
hgs2@hgs2-vm:~$
```

总共运行了5次test程序，统计如下表。

| 系统缺页次数(pfcount) | 当前进程缺页次数(pf) |
|-----------------|--------------|
| 629229 | 66 |
| 638594 | 68 |
| 639550 | 66 |
| 642137 | 67 |
| 642577 | 66 |

可以看出每次运行，系统缺页次数都会明显增长，而当前进程的缺页次数则基本不变。